

# AI in the Sciences and Engineering

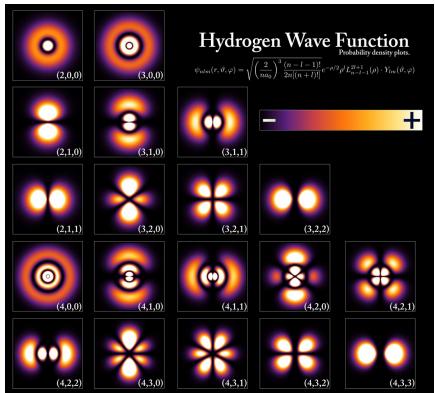
## Introduction to Physics- Informed Neural Networks

Spring Semester 2024

Siddhartha Mishra  
Ben Moseley

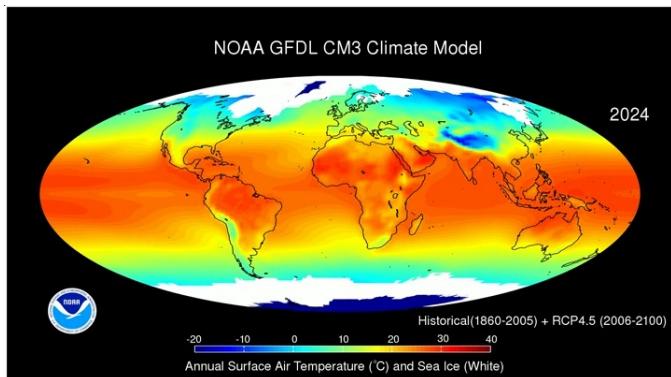
**ETH** zürich

# Recap – importance of PDEs



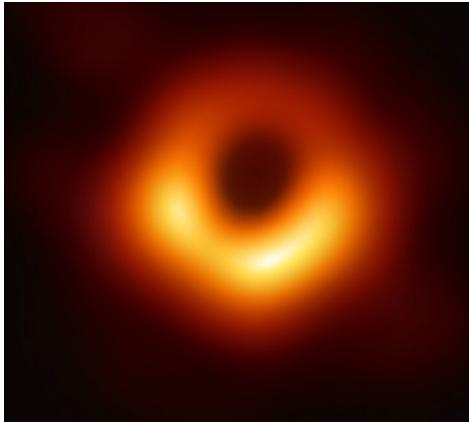
Source: Wikipedia

## Schrödinger equation



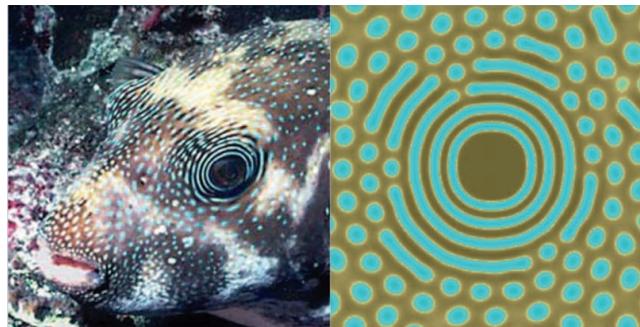
Source: NOAA

## Navier-Stokes equations



Source: The Event Horizon Telescope (2019)

## Einstein field equations



Source: Kondo and Miura, Science (2010)

## Reaction-diffusion equation

# Course timeline

Tutorials	Lectures
<i>Mon 12:15-14:00 HG E 5</i>	<i>Wed 08:15-10:00 ML H 44</i>
19.02.	21.02. <b>Course introduction</b>
26.02. Introduction to PyTorch	28.02. <del>Introduction to deep learning I</del>
04.03. CNNs and surrogate modelling	06.03. <b>Introduction to physics-informed neural networks</b>
11.03. Implementing PINNs I	13.03. PINNs – extensions and theory
18.03. Implementing PINNs II	20.03. Introduction to operator learning
25.03. Operator learning I	27.03. Fourier- and convolutional- neural operators
01.04.	03.04.
08.04. Operator learning II	10.04. Operator learning – limitations and extensions
15.04.	17.04. Foundational models for operator learning
22.04. GNNs	24.04. GNNs for PDEs / introduction to diffusion models
29.04. Transformers	01.05.
06.05. Diffusion models	08.05. Introduction to differentiable physics
13.05. Coding autodiff from scratch	15.05. Neural differential equations
20.05.	22.05. Symbolic regression and equation discovery
27.05. Introduction to JAX / NDEs	29.05. Guest lecture: ML in chemistry and biology
	<i>Fri 12:15-13:00 ML H 44</i>
	23.02. <del>Introduction to deep learning I</del>
	01.03. <del>Importance of PDEs in science</del>
	08.03. PINNs – limitations and extensions
	15.03. PINNs – theory
	22.03. DeepONets and spectral neural operators
	29.03.
	05.04.
	12.04. Introduction to transformers
	19.04. Graph neural networks for PDEs
	26.04. Introduction to diffusion models
	03.05. Diffusion models - applications
	10.05. Hybrid workflows
	17.05. Introduction to JAX
	24.05. Course summary and future trends
	31.05. Guest lecture: ML in chemistry and biology

# Lecture overview

- What are physics-informed neural networks (PINNs)?
- How to train PINNs
  - Live-coding a PINN in PyTorch
- Applications of PINNs
  - Simulation
  - Inversion
  - Equation discovery

# Lecture overview

- What are physics-informed neural networks (PINNs)?
- How to train PINNs
  - Live-coding a PINN in PyTorch
- Applications of PINNs
  - Simulation
  - Inversion
  - Equation discovery

# Learning objectives

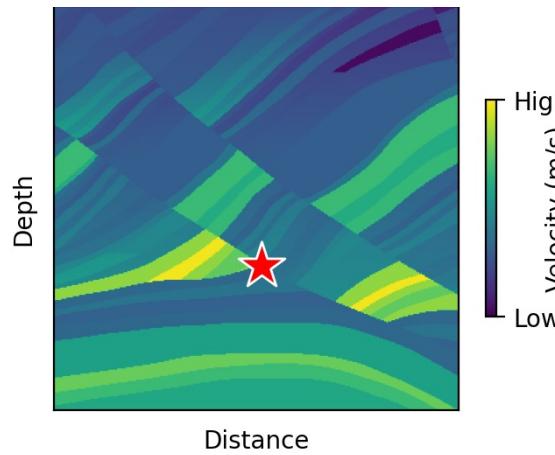
- Explain what a PINN is, and how they are trained
- Understand how to apply PINNs to different scientific tasks

# What is a physics-informed neural network (PINN)?

# Using neural networks for simulation



How could we solve this PDE using neural networks?



Wave equation:

$$\nabla^2 u - \frac{1}{c(x,y)^2} \frac{\partial^2 u}{\partial t^2} = s(x,y,t)$$

Initial conditions:

$$u(x, y, t = 0) = 0$$
$$u_t(x, y, t = 0) = 0$$

# What is a PINN?

Damped harmonic oscillator:

$$m \frac{d^2u}{dt^2} + \mu \frac{du}{dt} + ku = 0$$

Initial conditions:

$$\begin{aligned} u(t = 0) &= 1 \\ u_t(t = 0) &= 0 \end{aligned}$$

$u$  = displacement

$m$  = mass of oscillator

$\mu$  = coefficient of friction

$k$  = spring constant



Raissi et al, Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations, JCP (2018)  
Lagaris et al, Artificial neural networks for solving ordinary and partial differential equations, IEEE (1998)

# What is a PINN?

Damped harmonic oscillator:

$$m \frac{d^2u}{dt^2} + \mu \frac{du}{dt} + ku = 0$$

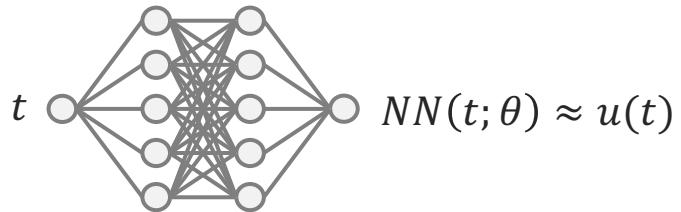
Key idea: use a neural network to directly approximate the solution



$$NN(t; \theta) \approx u(t)$$

Initial conditions:

$$\begin{aligned} u(t = 0) &= 1 \\ u_t(t = 0) &= 0 \end{aligned}$$



$u$  = displacement

$m$  = mass of oscillator

$\mu$  = coefficient of friction

$k$  = spring constant



Raissi et al, Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations, JCP (2018)  
Lagaris et al, Artificial neural networks for solving ordinary and partial differential equations, IEEE (1998)

# What is a PINN?

Damped harmonic oscillator:

$$m \frac{d^2u}{dt^2} + \mu \frac{du}{dt} + ku = 0$$

Initial conditions:

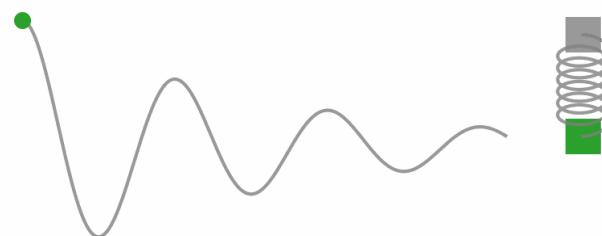
$$\begin{aligned} u(t=0) &= 1 \\ u_t(t=0) &= 0 \end{aligned}$$

$u$  = displacement

$m$  = mass of oscillator

$\mu$  = coefficient of friction

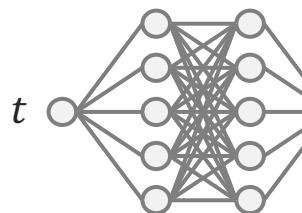
$k$  = spring constant



Key idea: use a neural network to directly approximate the solution



$$NN(t; \theta) \approx u(t)$$



Train the network using the loss function:

Boundary loss

$$L_b(\theta)$$

Physics loss

$$L_p(\theta)$$

(aka PDE residual)

$$\left. \begin{aligned} L(\theta) &= \lambda_1(NN(t=0; \theta) - 1)^2 \\ &+ \lambda_2 \left( \frac{dNN}{dt}(t=0; \theta) - 0 \right)^2 \\ &+ \frac{1}{N_p} \sum_i^{N_p} \left( \left[ m \frac{d^2}{dt^2} + \mu \frac{d}{dt} + k \right] NN(t_i; \theta) \right)^2 \end{aligned} \right\}$$

Raissi et al, Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations, JCP (2018)  
Lagaris et al, Artificial neural networks for solving ordinary and partial differential equations, IEEE (1998)

# What is a PINN?

Damped harmonic oscillator:

$$m \frac{d^2u}{dt^2} + \mu \frac{du}{dt} + ku = 0$$

Initial conditions:

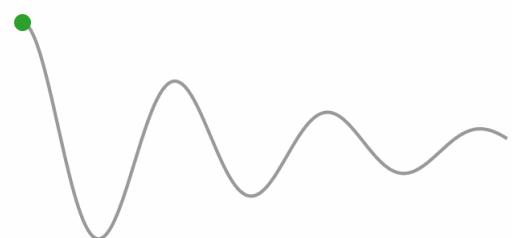
$$\begin{aligned} u(t=0) &= 1 \\ u_t(t=0) &= 0 \end{aligned}$$

$u$  = displacement

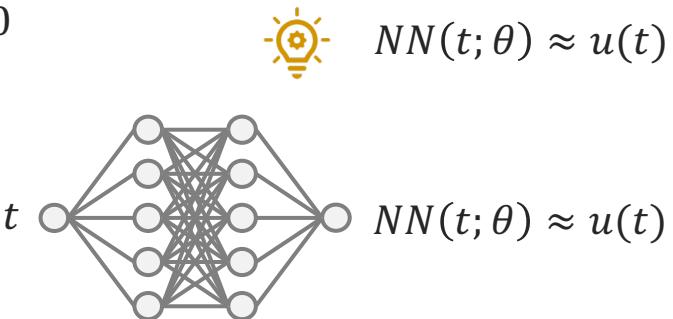
$m$  = mass of oscillator

$\mu$  = coefficient of friction

$k$  = spring constant

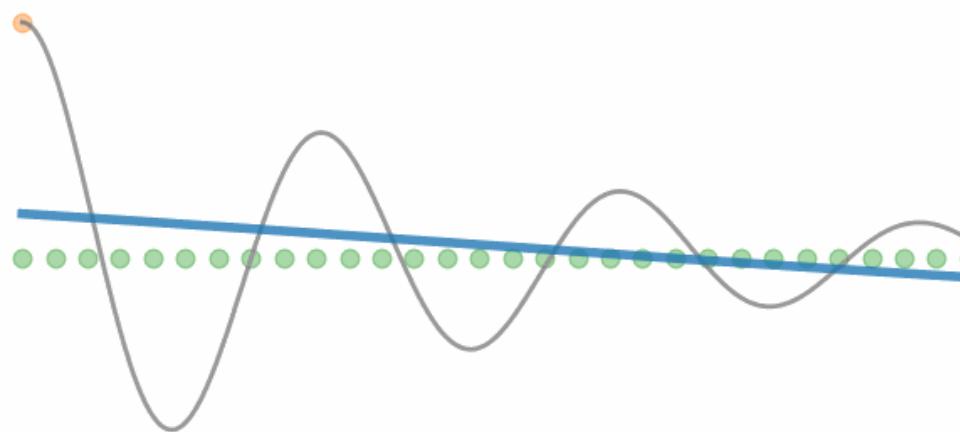


Key idea: use a neural network to directly approximate the solution



Train the network using the loss function:

$$\begin{aligned} \text{Boundary loss } L_b(\theta) &= \left[ L(\theta) = \lambda_1 (NN(t=0; \theta) - 1)^2 \right. \\ &\quad \left. + \lambda_2 \left( \frac{dNN}{dt}(t=0; \theta) - 0 \right)^2 \right] \\ \text{Physics loss } L_p(\theta) &= \left[ \right. \\ &\quad \left. \left( \text{aka PDE residual} \right) \right. \\ &\quad \left. + \frac{1}{N_p} \sum_i^{N_p} \left( \left[ m \frac{d^2}{dt^2} + \mu \frac{d}{dt} + k \right] NN(t_i; \theta) \right)^2 \right] \end{aligned}$$



Training step: 150

- Exact solution
- Neural network prediction
- Boundary loss training locations
- Physics loss training locations

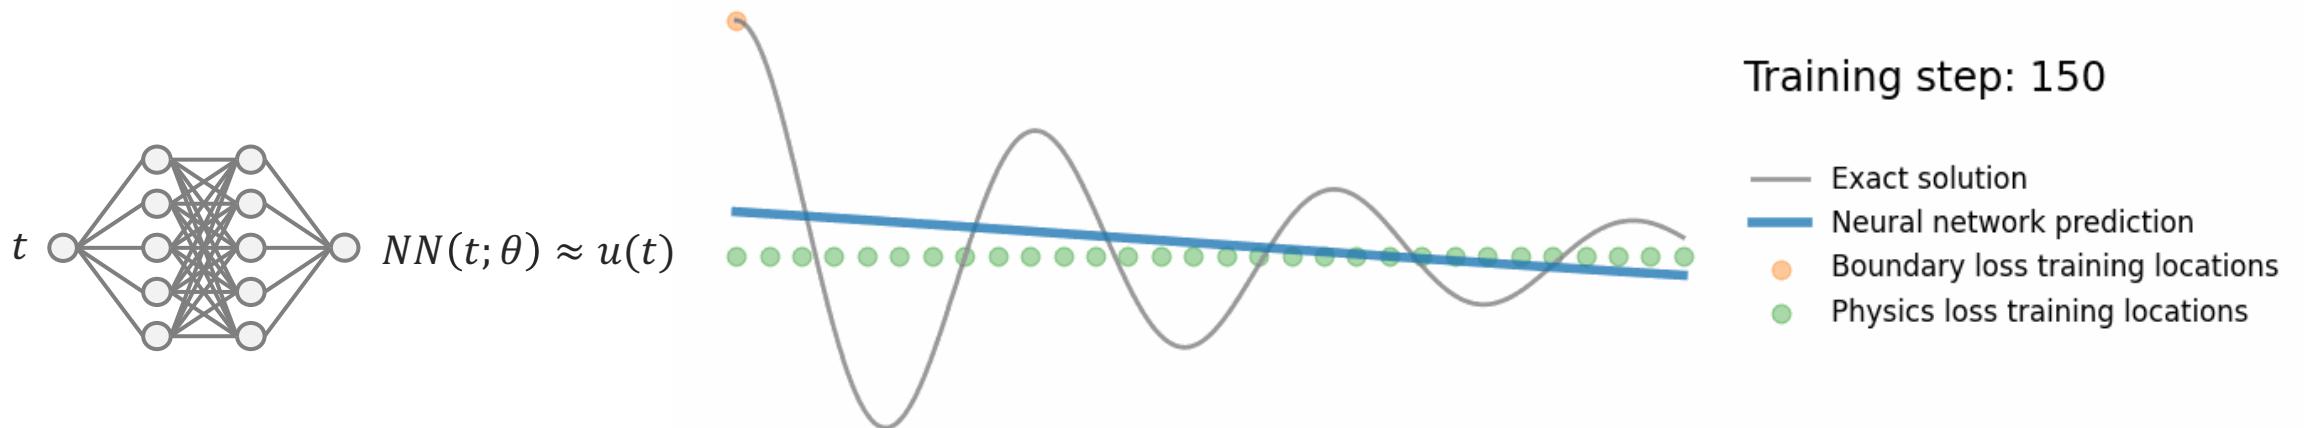
Raissi et al, Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations, JCP (2018)  
Lagaris et al, Artificial neural networks for solving ordinary and partial differential equations, IEEE (1998)

# What is a PINN?

- $\{t_i\}_{i=1}^{N_p}$  are known as **collocation points**, which are sampled throughout the domain
- $\lambda_1, \lambda_2$  are scalar hyperparameters which **balance** the contribution of each loss term

Train the network using the loss function:

$$\begin{aligned} \text{Boundary loss } L_b(\theta) & \left[ \begin{array}{l} L(\theta) = \lambda_1 (NN(t=0; \theta) - 1)^2 \\ + \lambda_2 \left( \frac{dNN}{dt}(t=0; \theta) - 0 \right)^2 \end{array} \right] \\ \text{Physics loss } L_p(\theta) & \left[ \begin{array}{l} \text{(aka PDE residual)} \\ + \frac{1}{N_p} \sum_i^{N_p} \left( \left[ m \frac{d^2}{dt^2} + \mu \frac{d}{dt} + k \right] NN(t_i; \theta) \right)^2 \end{array} \right] \end{aligned}$$



Raissi et al, Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations, JCP (2018)  
Lagaris et al, Artificial neural networks for solving ordinary and partial differential equations, IEEE (1998)

# PINNs are a general framework for solving PDEs

Given a PDE and its boundary/initial conditions

$$\begin{aligned}\mathcal{D}[u(x)] &= f(x), \quad x \in \Omega \subset \mathbb{R}^d \\ \mathcal{B}_k[u(x)] &= g_k(x), \quad x \in \Gamma_k \subset \partial\Omega\end{aligned}$$

Where  $\mathcal{D}$  is some differential operator,  $\mathcal{B}_k$  are a set of boundary operators, and  $u(x)$  is the solution to the PDE

PINNs train a neural network to **approximate** the solution to the PDE  $NN(x; \theta) \approx u(x)$  using the following loss function:

$$L(\theta) = L_b(\theta) + L_p(\theta)$$

$$L_b(\theta) = \sum_k \frac{\lambda_k}{N_{bk}} \sum_j^{|N_{bk}|} \| \mathcal{B}_k[NN(x_{kj}; \theta)] - g_k(x_{kj}) \|^2 \text{ Boundary loss}$$

$$L_p(\theta) = \frac{1}{N_p} \sum_i^{N_p} \| \mathcal{D}[NN(x_i; \theta)] - f(x_i) \|^2 \text{ Physics loss}$$

# PINNs are a general framework for solving PDEs

Given a PDE and its boundary/initial conditions

$$\begin{aligned}\mathcal{D}[u(x)] &= f(x), \quad x \in \Omega \subset \mathbb{R}^d \\ \mathcal{B}_k[u(x)] &= g_k(x), \quad x \in \Gamma_k \subset \partial\Omega\end{aligned}$$

Where  $\mathcal{D}$  is some differential operator,  $\mathcal{B}_k$  are a set of boundary operators, and  $u(x)$  is the solution to the PDE

PINNs train a neural network to **approximate** the solution to the PDE  $NN(x; \theta) \approx u(x)$  using the following loss function:

$$L(\theta) = L_b(\theta) + L_p(\theta)$$

$$L_b(\theta) = \sum_k \frac{\lambda_k}{N_{bk}} \sum_j^{N_{bk}} \| \mathcal{B}_k[NN(x_{kj}; \theta)] - g_k(x_{kj}) \|^2 \text{ Boundary loss}$$

$$L_p(\theta) = \frac{1}{N_p} \sum_i^{N_p} \| \mathcal{D}[NN(x_i; \theta)] - f(x_i) \|^2 \text{ Physics loss}$$

For example, the 1+1D viscous Burgers' equation:

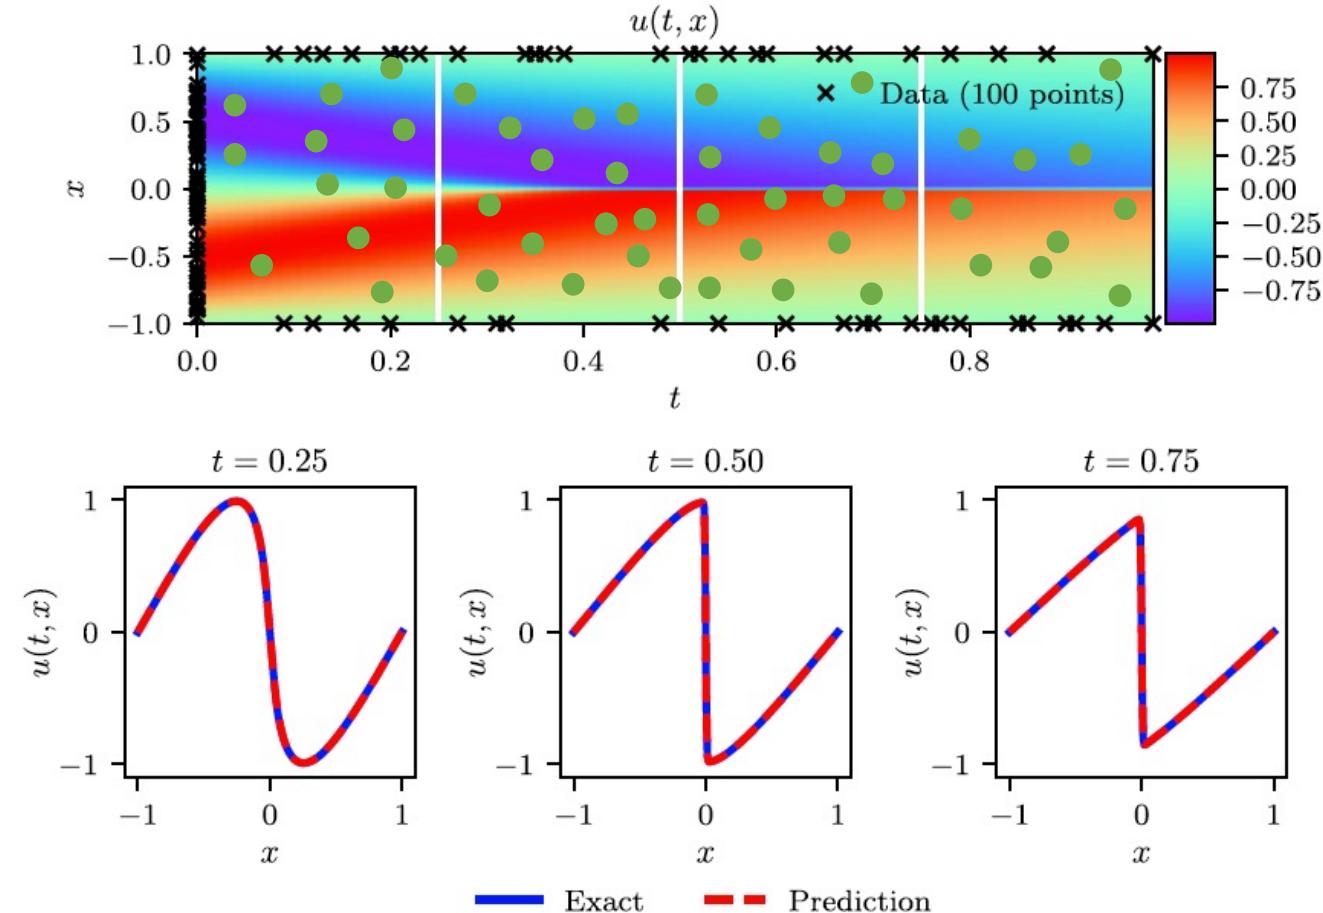
$$\frac{\partial u}{\partial t} + u \frac{\partial u}{\partial x} - \nu \frac{\partial^2 u}{\partial x^2} = 0$$

$$\begin{aligned}u(x, 0) &= -\sin(\pi x) \\ u(-1, t) &= u(+1, t) = 0\end{aligned}$$

$$u(x, t) \approx NN(x, t; \theta)$$

$$\begin{aligned}L_b(\theta) &= \frac{\lambda_1}{N_{b1}} \sum_j^{N_{b1}} (NN(x_j, 0; \theta) + \sin(\pi x_j))^2 \\ &\quad + \frac{\lambda_2}{N_{b2}} \sum_k^{N_{b2}} (NN(-1, t_k; \theta) - 0)^2 \\ &\quad + \frac{\lambda_3}{N_{b3}} \sum_l^{N_{b3}} (NN(+1, t_l; \theta) - 0)^2 \\ L_p(\theta) &= \frac{1}{N_p} \sum_i^{N_p} \left( \left( \frac{\partial NN}{\partial t} + NN \frac{\partial NN}{\partial x} - \nu \frac{\partial^2 NN}{\partial x^2} \right)(x_i, t_i; \theta) \right)^2\end{aligned}$$

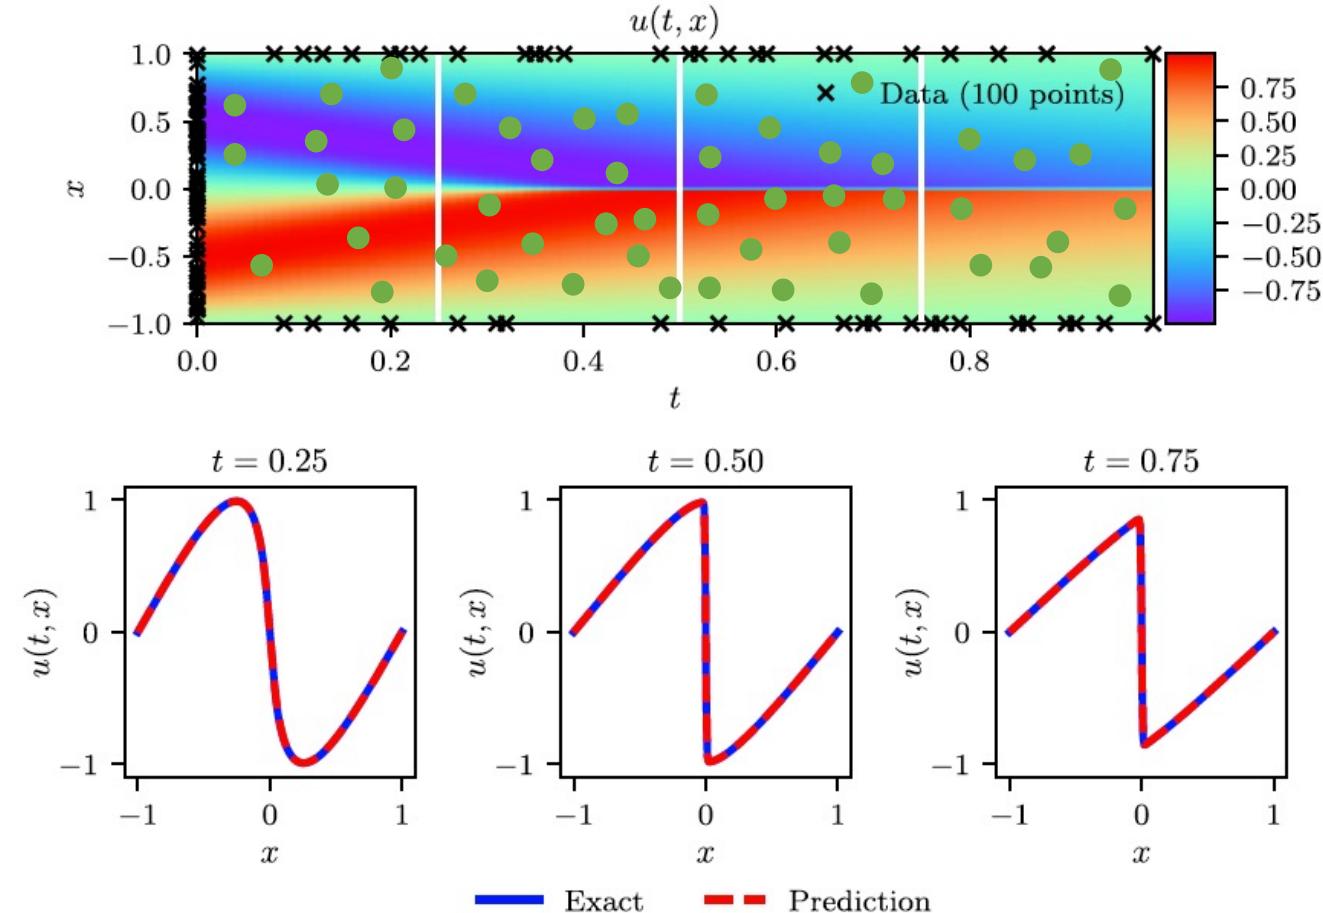
# PINNs for solving viscous Burgers' equation



$$L_b(\theta) = \frac{\lambda_1}{N_{b1}} \sum_j^{N_{b1}} (NN(x_j, 0; \theta) + \sin(\pi x_j))^2$$
$$+ \frac{\lambda_2}{N_{b2}} \sum_k^{N_{b2}} (NN(-1, t_k; \theta) - 0)^2$$
$$+ \frac{\lambda_3}{N_{b3}} \sum_l^{N_{b3}} (NN(+1, t_l; \theta) - 0)^2$$
$$L_p(\theta) = \frac{1}{N_p} \sum_i^{N_p} \left( \left( \frac{\partial NN}{\partial t} + NN \frac{\partial NN}{\partial x} - \nu \frac{\partial^2 NN}{\partial x^2} \right) (x_i, t_i; \theta) \right)^2$$

Raissi et al, Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations, JCP (2018)

# PINNs for solving viscous Burgers' equation



Raissi et al, Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations, JCP (2018)

$$L_b(\theta) = \frac{\lambda_1}{N_{b1}} \sum_j^{N_{b1}} (NN(x_j, 0; \theta) + \sin(\pi x_j))^2$$
$$+ \frac{\lambda_2}{N_{b2}} \sum_k^{N_{b2}} (NN(-1, t_k; \theta) - 0)^2$$
$$+ \frac{\lambda_3}{N_{b3}} \sum_l^{N_{b3}} (NN(+1, t_l; \theta) - 0)^2$$
$$L_p(\theta) = \frac{1}{N_p} \sum_i^{N_p} \left( \left( \frac{\partial NN}{\partial t} + NN \frac{\partial NN}{\partial x} - \nu \frac{\partial^2 NN}{\partial x^2} \right) (x_i, t_i; \theta) \right)^2$$

$$\nu = 0.01/\pi$$

$N_p = 10,000$  (Latin hypercube sampling)

$$N_{b1} + N_{b2} + N_{b3} = 100$$

Fully connected network with 9 layers, 20 hidden units (3021 free parameters)

Tanh activation function

L-BFGS optimiser

# PINNs – an entire research field

## SPRINGER LINK

Find a journal   Publish with us   Track your research

Search

Home > Journal of Scientific Computing > Article

## Scientific Machine Learning Through Physics-Informed Neural Networks: Where we are and What's Next

Open access | Published: 26 July 2022

Volume 92, article number 88, (2022) Cite this article

Download PDF

You have full access to this open access article

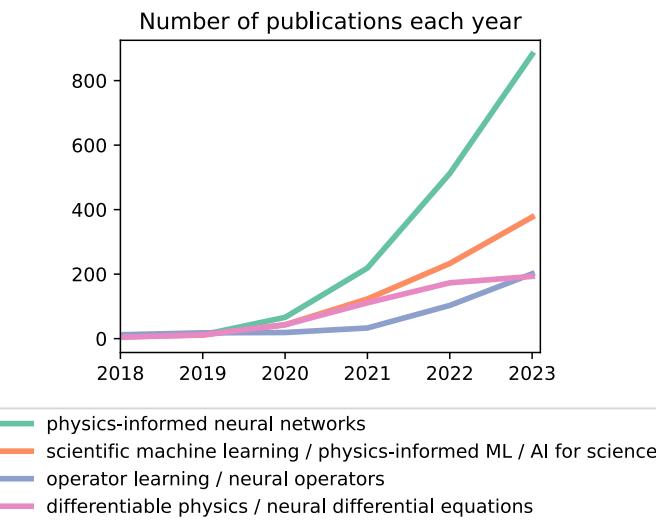
Salvatore Cuomo, Vincenzo Schiano Di Cola, Fabio Giampaolo, Gianluigi Rozza, Maziar Raissi & Francesco Piccialli

67k Accesses 274 Citations 7 Altmetric Explore all metrics →

### Abstract

Physics-Informed Neural Networks (PINN) are neural networks (NNs) that encode model equations, like Partial Differential Equations (PDE), as a component of the neural network itself. PINNs are nowadays used to solve PDEs, fractional equations, integral-differential equations, and stochastic PDEs. This novel methodology has arisen as a multi-task learning framework in which a NN must fit observed data while reducing a PDE residual.

This article provides a comprehensive review of the literature on PINNs: while the primary goal of the studies was to characterize these networks and their related advantages and



Source: Scopus keyword search (Feb 2024)

- The basic concepts behind PINNs were introduced in the 1990s (Lagaris et al, IEEE (1998) and others)
- Raissi et al, JCP (2018) is the seminal paper which reimplemented and extended PINNs using a modern deep learning framework

Raissi et al, Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations, JCP (2018)  
Lagaris et al, Artificial neural networks for solving ordinary and partial differential equations, IEEE (1998)

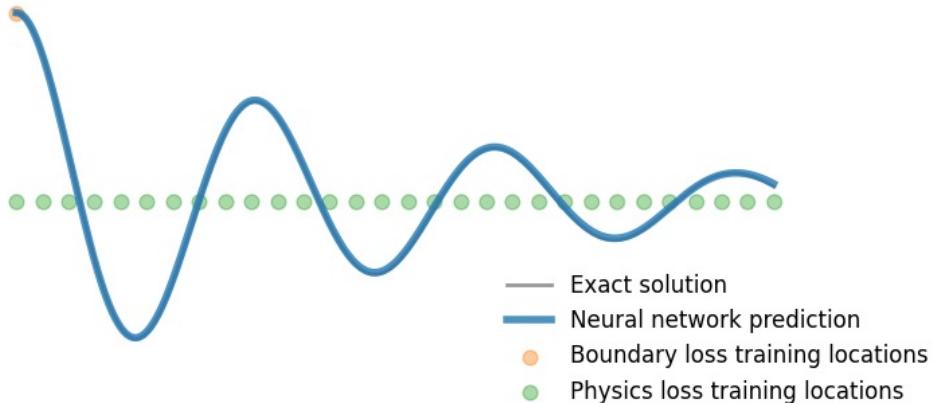
# Lecture overview

- What are physics-informed neural networks (PINNs)?
- How to train PINNs
  - Live-coding a PINN in PyTorch
- Applications of PINNs
  - Simulation
  - Inversion
  - Equation discovery

# Learning objectives

- Explain what a PINN is, and how they are trained
- Understand how to apply PINNs to different scientific tasks

# PINN training loop

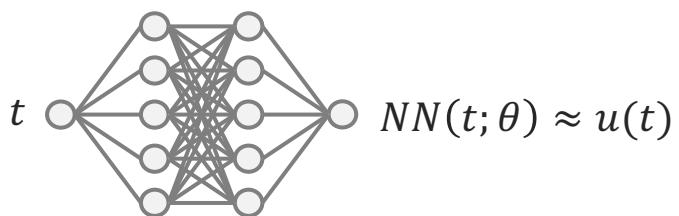


Boundary loss  $L_b(\theta)$

$$\left\{ \begin{array}{l} L(\theta) = \lambda_1(NN(t=0; \theta) - 1)^2 \\ \quad + \lambda_2 \left( \frac{dNN}{dt}(t=0; \theta) - 0 \right)^2 \end{array} \right.$$

Physics loss  $L_p(\theta)$

$$\left\{ \begin{array}{l} \quad + \frac{1}{N_p} \sum_i^{N_p} \left( \left[ m \frac{d^2}{dt^2} + \mu \frac{d}{dt} + k \right] NN(t_i; \theta) \right)^2 \end{array} \right.$$



Training loop:

1. Sample boundary/ physics training points
2. Compute network outputs
3. Compute 1<sup>st</sup> and 2<sup>nd</sup> order gradient of network output **with respect to network input**
4. Compute loss
5. Compute gradient of loss function **with respect to network parameters**
6. Take gradient descent step

How can we compute the gradients (e.g.  $\frac{dNN}{dt}$  and  $\frac{dL}{d\theta}$ ) required in steps 3 and 5?

# Gradient computation for PINNs

Assume network is a MLP, e.g.:

$$NN(t; \theta) = W_3(\sigma(W_2\sigma(W_1t + \mathbf{b}_1) + \mathbf{b}_2) + b_3) \approx u(t)$$

Analogous to backpropagation, we can use the multivariate **chain rule** to compute gradients with respect to network **inputs**:

$$NN(t; \theta) = f \circ \mathbf{g} \circ \mathbf{h}(t; \theta)$$

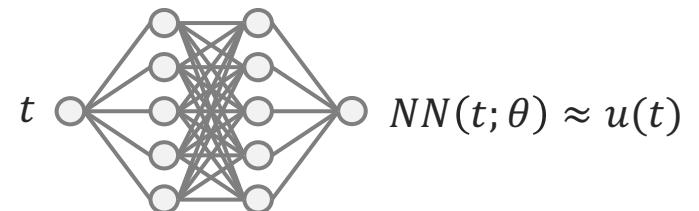
$$f = W_3\sigma(\mathbf{g}) + b_3$$

$$\mathbf{g} = W_2\sigma(\mathbf{h}) + \mathbf{b}_2$$

$$\mathbf{h} = W_1t + \mathbf{b}_1$$

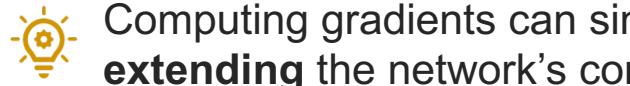
$$\frac{\partial NN}{\partial t} = \frac{\partial f}{\partial \mathbf{g}} \frac{\partial \mathbf{g}}{\partial \mathbf{h}} \frac{\partial \mathbf{h}}{\partial t}$$

$$\frac{\partial NN}{\partial t} = W_3 \text{ diag}(\sigma'(\mathbf{g})) W_2 \text{ diag}(\sigma'(\mathbf{h})) W_1$$



(see Lecture 2: Introduction to Deep Learning Part I)

# Extending computational graph



Computing gradients can simply be thought of as **extending** the network's computational graph:

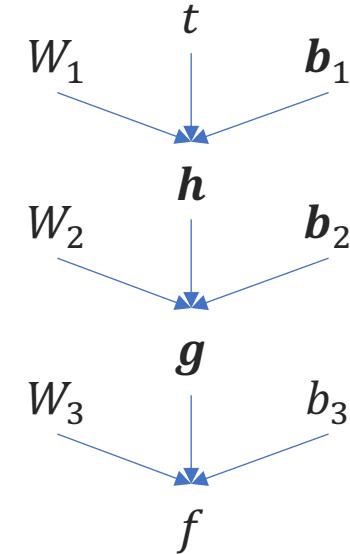
$$NN(t; \theta) = W_3(\sigma(W_2\sigma(W_1t + \mathbf{b}_1) + \mathbf{b}_2) + b_3)$$

$$f = W_3\sigma(\mathbf{g}) + b_3$$

$$\mathbf{g} = W_2\sigma(\mathbf{h}) + \mathbf{b}_2$$

$$\mathbf{h} = W_1t + \mathbf{b}_1$$

$$\frac{\partial NN}{\partial t} = W_3 \text{ diag}(\sigma'(\mathbf{g}))W_2 \text{ diag}(\sigma'(\mathbf{h}))W_1$$



# Extending computational graph



Computing gradients can simply be thought of as **extending** the network's computational graph:

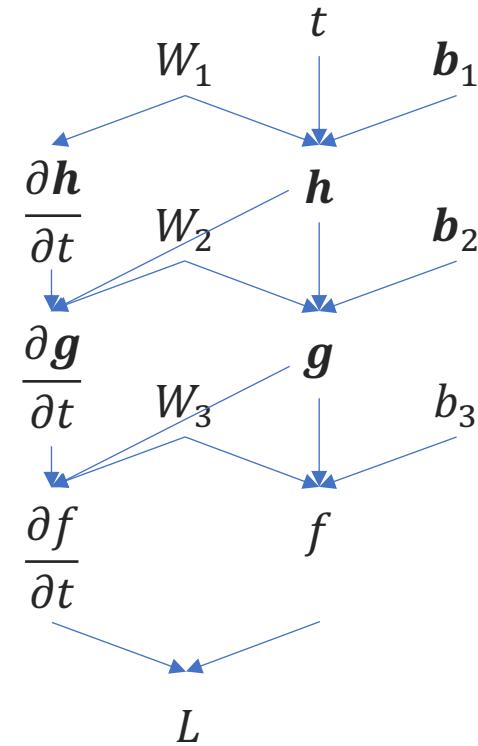
$$NN(t; \theta) = W_3(\sigma(W_2\sigma(W_1t + \mathbf{b}_1) + \mathbf{b}_2) + b_3)$$

$$f = W_3\sigma(\mathbf{g}) + b_3$$

$$\mathbf{g} = W_2\sigma(\mathbf{h}) + \mathbf{b}_2$$

$$\mathbf{h} = W_1t + \mathbf{b}_1$$

$$\frac{\partial NN}{\partial t} = W_3 \text{ diag}(\sigma'(\mathbf{g}))W_2 \text{ diag}(\sigma'(\mathbf{h}))W_1$$



# Extending computational graph

💡 Computing gradients can simply be thought of as **extending** the network's computational graph:

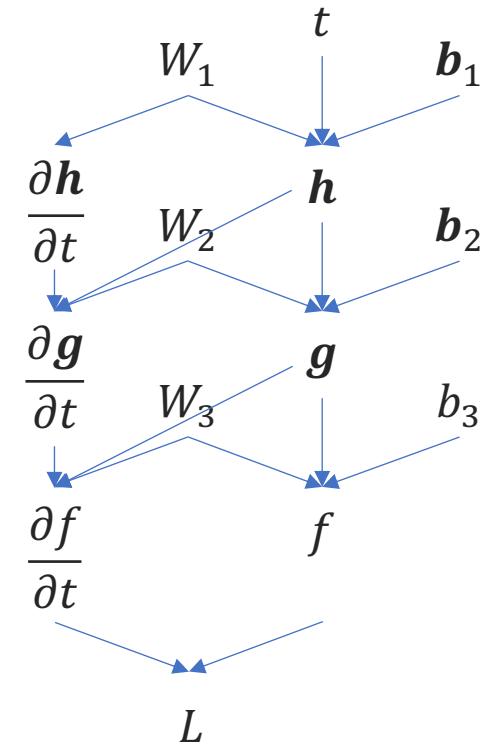
$$NN(t; \theta) = W_3(\sigma(W_2\sigma(W_1t + \mathbf{b}_1) + \mathbf{b}_2) + b_3)$$

$$f = W_3\sigma(\mathbf{g}) + b_3$$

$$\mathbf{g} = W_2\sigma(\mathbf{h}) + \mathbf{b}_2$$

$$\mathbf{h} = W_1t + \mathbf{b}_1$$

$$\frac{\partial NN}{\partial t} = W_3 \text{ diag}(\sigma'(\mathbf{g}))W_2 \text{ diag}(\sigma'(\mathbf{h}))W_1$$



💡 We can recursively apply **autodifferentiation** to compute gradients and extend the graph

# PINN training loop

Training loop:

1. Sample boundary/ physics training points
2. Compute network outputs
3. Compute 1<sup>st</sup> and 2<sup>nd</sup> order gradient of network output **with respect to network input <= (recursively) apply autodiff, extending graph**
4. Compute loss
5. Compute gradient of loss function **with respect to network parameters <= apply autodiff on extended graph**
6. Take gradient descent step

```
# PINN training psuedocode

#2.
t.requires_grad_(True)# tells PyTorch to start tracking graph
theta.requires_grad_(True)
u = NN(t, theta)

#3.
dudt = torch.autograd.grad(u, t, torch.ones_like(u),
                           create_graph=True)[0]
d2udt2 = torch.autograd.grad(dudt, t, torch.ones_like(u),
                             create_graph=True)[0]

#4.
physics_loss = torch.mean((m*d2udt2 + mu*dudt + k*u)**2)
loss = physics_loss + lambda_*boundary_loss

#5.
dtheta = torch.autograd.grad(loss, theta)[0]
```

 We can recursively apply **autodifferentiation** to compute gradients and extend the graph

# Live-coding a PINN in PyTorch

# Computational cost of higher order derivatives



Note, gradient computation roughly **doubles** the size of the computational graph\*

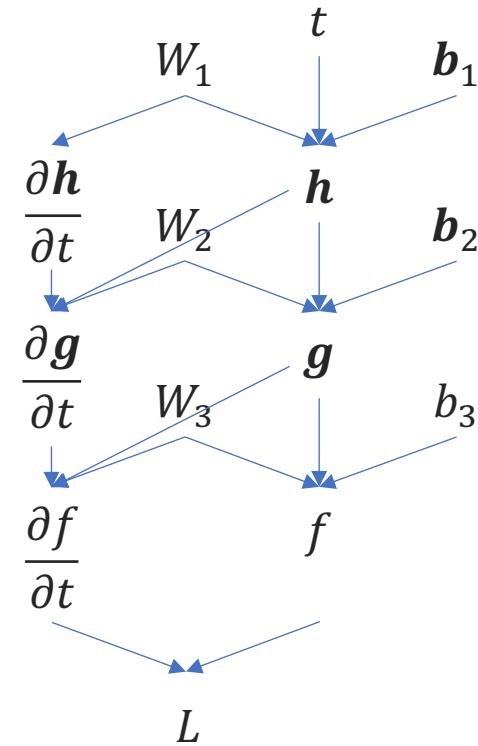
- ⇒ The cost of evaluating  $\frac{\partial^n NN}{\partial t^n}$  grows **exponentially** with  $n$  (!)
- ⇒ Most time is spent computing gradients, not the forward pass, when training PINNs

\*More precisely, given some  $f: \mathbb{R}^n \rightarrow \mathbb{R}^m$ , its Jacobian  $J \in \mathbb{R}^{m \times n}$  and some vector  $v \in \mathbb{R}^n$

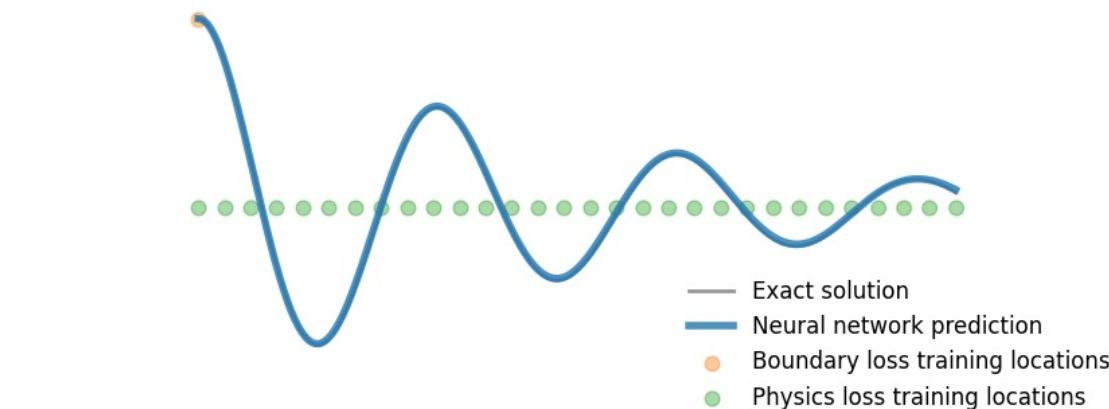
$$TIME\{f, Jv\} \leq \omega TIME\{f\}$$

With a constant  $\omega \in [2, 2.5]$  using autodifferentiation

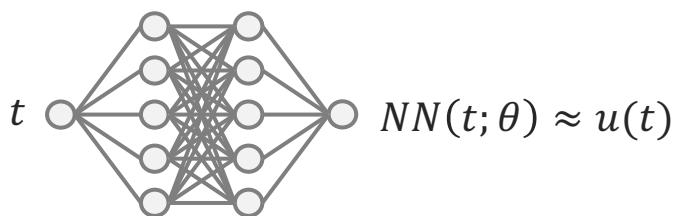
For detailed derivation, see eg: Griewank and Walther, Evaluating Derivatives, Ch 3.1, SIAM (2008))



# Other important training considerations

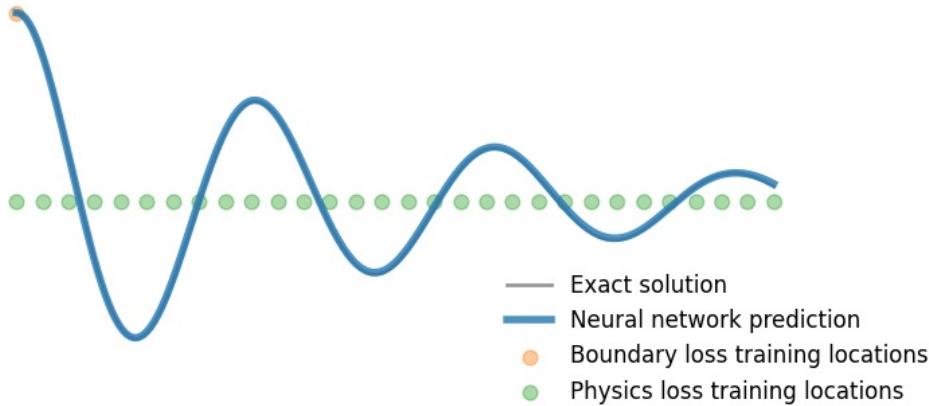


$$\begin{aligned} \text{Boundary loss } L_b(\theta) &= \left\{ \begin{array}{l} L(\theta) = \lambda_1(NN(t=0; \theta) - 1)^2 \\ \quad + \lambda_2 \left( \frac{dNN}{dt}(t=0; \theta) - 0 \right)^2 \end{array} \right. \\ \text{Physics loss } L_p(\theta) &= \left\{ \begin{array}{l} \quad + \frac{1}{N_p} \sum_i^{N_p} \left( \left[ m \frac{d^2}{dt^2} + \mu \frac{d}{dt} + k \right] NN(t_i; \theta) \right)^2 \end{array} \right. \end{aligned}$$

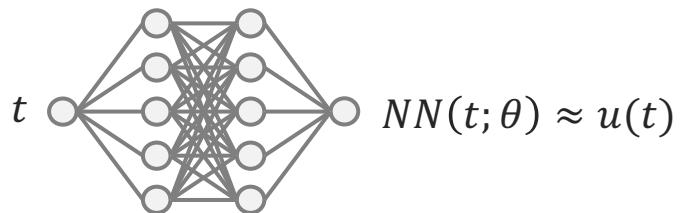


- $\lambda$  values often significantly affect convergence
- “Enough” collocation points  $\{t_i\}_{i=1}^{N_p}$  must be chosen so that the learned solution is accurate across the full domain
- Similarly, “enough” boundary points must be chosen such that the learned solution is unique
- Training points are usually sampled uniformly, randomly or quasirandomly
- PINNs still suffer from approximation, estimation and optimisation error (just like normal neural networks, see Lecture 3: Introduction to Deep Learning Part II)!

# PINNs from a ML perspective



$$\begin{aligned} \text{Boundary loss } L_b(\theta) &= \lambda_1(NN(t=0; \theta) - 1)^2 \\ &\quad + \lambda_2 \left( \frac{dNN}{dt}(t=0; \theta) - 0 \right)^2 \\ \text{Physics loss } L_p(\theta) &= \frac{1}{N_p} \sum_i^{N_p} \left( \left[ m \frac{d^2}{dt^2} + \mu \frac{d}{dt} + k \right] NN(t_i; \theta) \right)^2 \end{aligned}$$

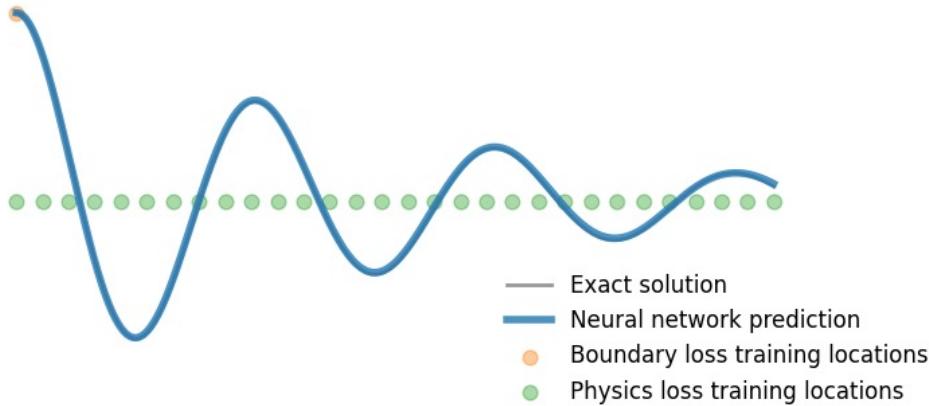


From a mathematical perspective:

- PINNs provide a way to solve PDEs:
  - Neural network is a mesh-free, **functional** approximation of PDE solution
  - Physics loss is used to assert solution is **consistent** with PDE
  - Boundary loss is used to assert boundary/initial conditions, to ensure solution is **unique**

What can we say about PINNs from a ML perspective?

# PINNs from a ML perspective



$$\begin{aligned} \text{Boundary loss } L_b(\theta) &= \left\{ \begin{array}{l} L(\theta) = \lambda_1(NN(t=0; \theta) - 1)^2 \\ \quad + \lambda_2 \left( \frac{dNN}{dt}(t=0; \theta) - 0 \right)^2 \end{array} \right. \\ \text{Physics loss } L_p(\theta) &= \left\{ \begin{array}{l} \quad + \frac{1}{N_p} \sum_i^{N_p} \left( \left[ m \frac{d^2}{dt^2} + \mu \frac{d}{dt} + k \right] NN(t_i; \theta) \right)^2 \end{array} \right. \end{aligned}$$

$t$   $NN(t; \theta) \approx u(t)$

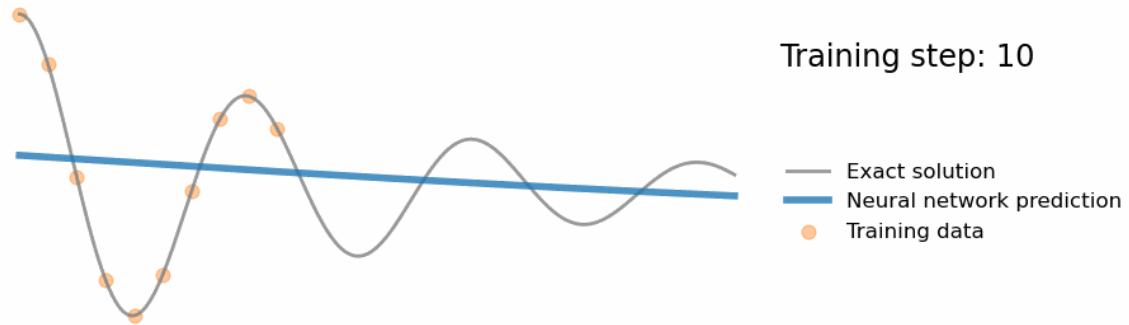
From a mathematical perspective:

- PINNs provide a way to solve PDEs:
  - Neural network is a mesh-free, **functional** approximation of PDE solution
  - Physics loss is used to assert solution is **consistent** with PDE
  - Boundary loss is used to assert boundary/initial conditions, to ensure solution is **unique**

From a ML perspective:

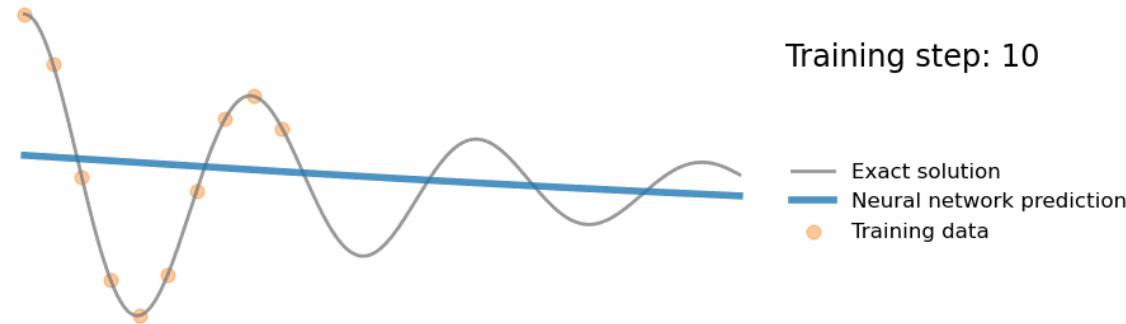
- Physics loss is an **unsupervised** regulariser, which  adds prior knowledge

# PINNs from a ML perspective

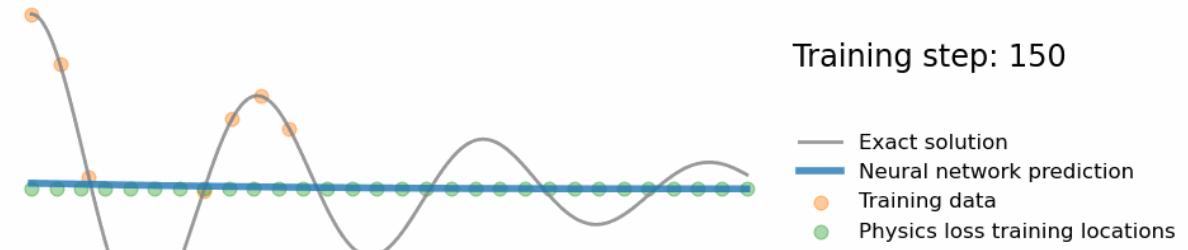


$$L(\theta) = \frac{1}{N} \sum_i^N (NN(t_i; \theta) - u_i)^2 \quad \text{Supervised loss}$$

# PINNs from a ML perspective



$$L(\theta) = \frac{1}{N} \sum_i^N (NN(t_i; \theta) - u_i)^2 \quad \text{Supervised loss}$$



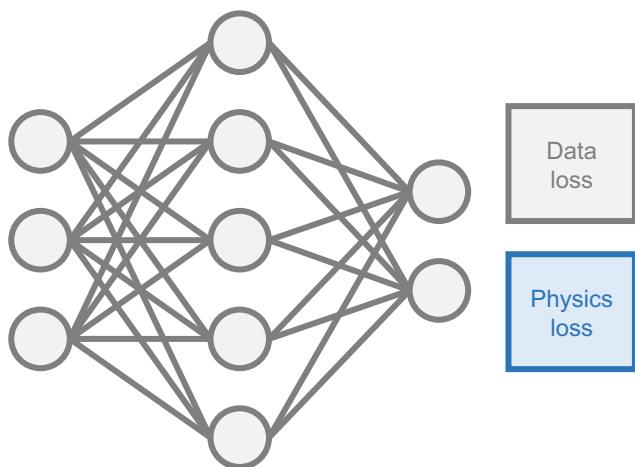
$$\begin{aligned} L(\theta) &= \frac{1}{N} \sum_i^N (NN(t_i; \theta) - u_i)^2 \quad \text{Supervised loss} \\ &+ \frac{\lambda}{M} \sum_j^M \left( \left[ m \frac{d^2}{dt^2} + \mu \frac{d}{dt} + k \right] NN(t_j; \theta) \right)^2 \quad \text{Physics loss} \end{aligned}$$

From a ML perspective:

- Physics loss is an **unsupervised** regulariser, which adds prior knowledge

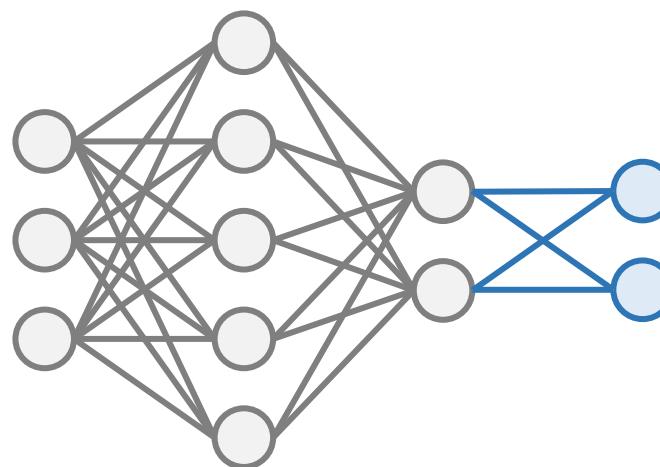
# Ways to incorporate scientific principles into machine learning

## Loss function



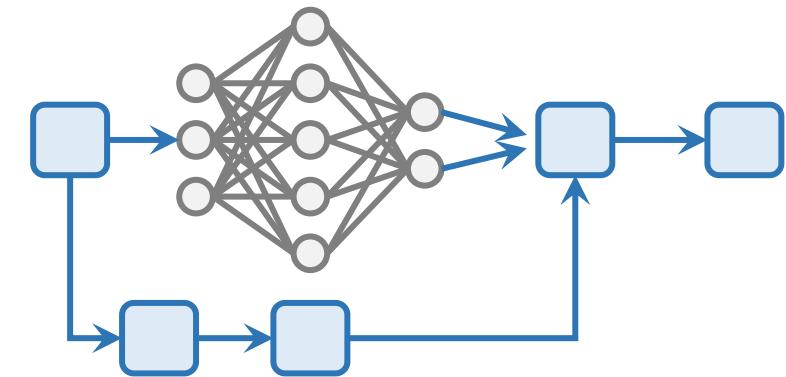
Example:  
Physics-informed neural networks  
(add governing equations to loss  
function)

## Architecture



Example:  
Encoding symmetries / conservation laws  
(e.g. energy conservation, rotational  
invariance)

## Hybrid approaches



Example:  
Neural differential equations  
(incorporating neural networks into PDE  
models)

# Lecture overview

- What are physics-informed neural networks (PINNs)?
- How to train PINNs
  - Live-coding a PINN in PyTorch
- Applications of PINNs
  - Simulation
  - Inversion
  - Equation discovery

# Learning objectives

- Explain what a PINN is, and how they are trained
- Understand how to apply PINNs to different scientific tasks

# 5 min break

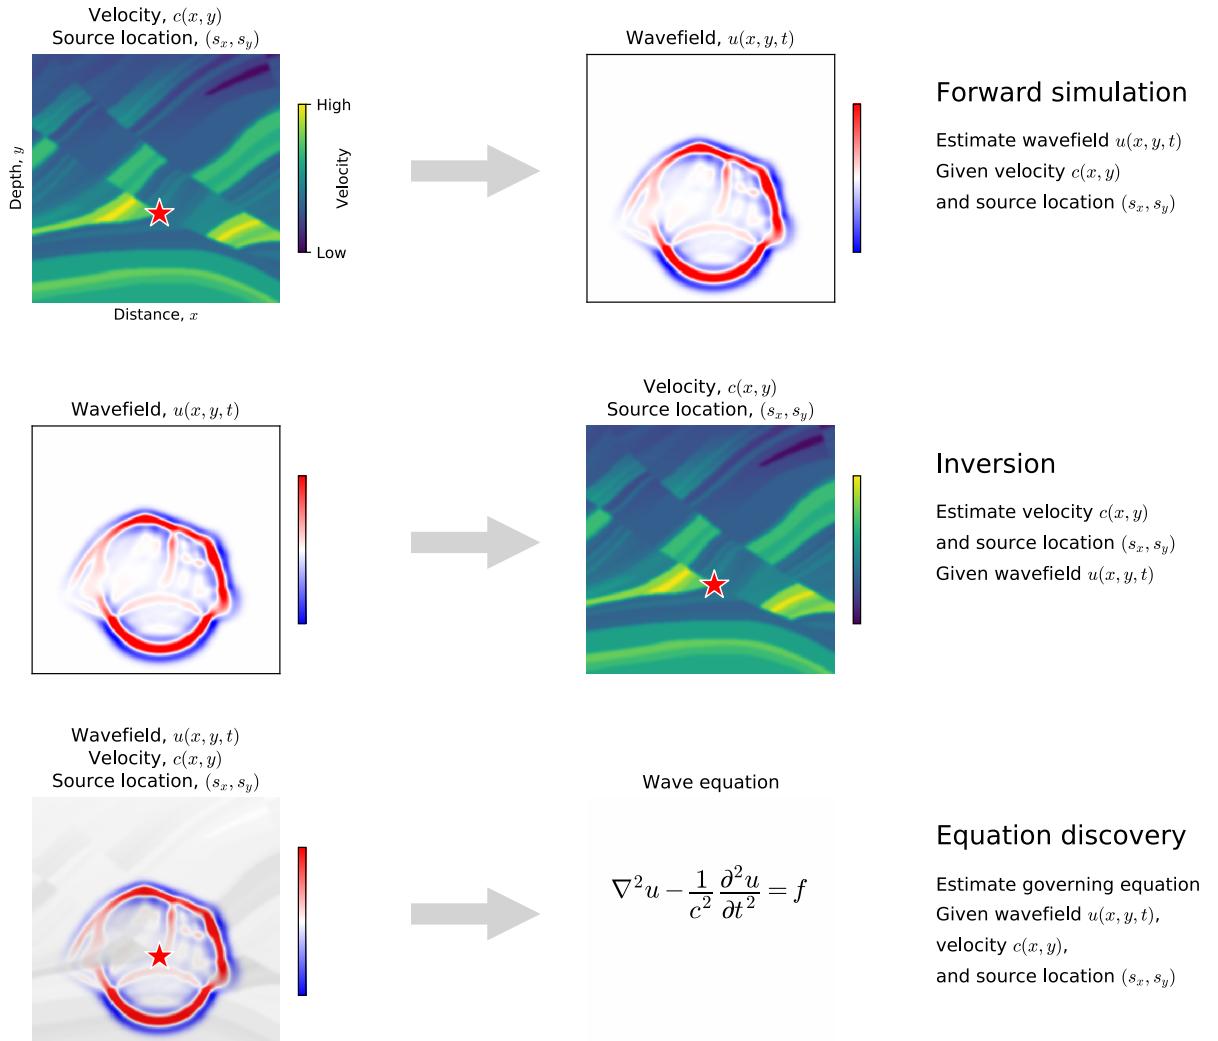
# Lecture overview

- What are physics-informed neural networks (PINNs)?
- How to train PINNs
  - Live-coding a PINN in PyTorch
- Applications of PINNs
  - Simulation
  - Inversion
  - Equation discovery

# Learning objectives

- Explain what a PINN is, and how they are trained
- Understand how to apply PINNs to different scientific tasks

# Key scientific tasks



PINNs can be used to solve **forward**, **inverse** and **equation discovery** problems related to PDEs

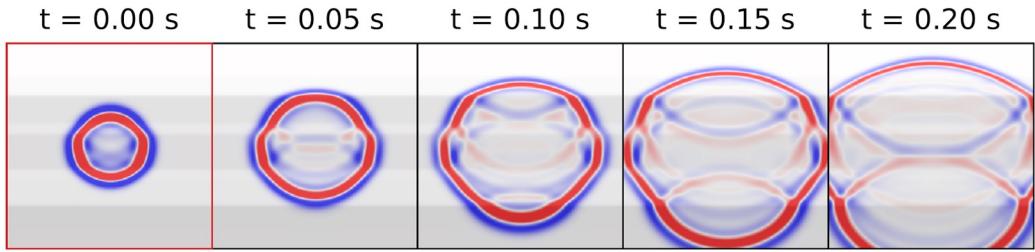
$a$  = set of input conditions

$F$  = physical model of the system (usually a PDE)

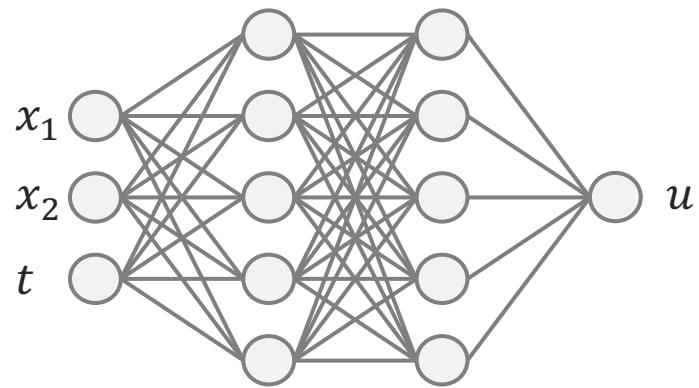
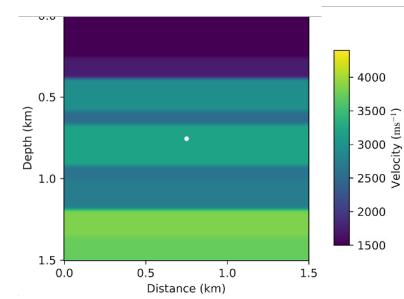
$b$  = resulting properties given  $F$  and  $a$

# PINNs for solving wave equation

Ground truth FD simulation



Velocity model,  $c(x)$



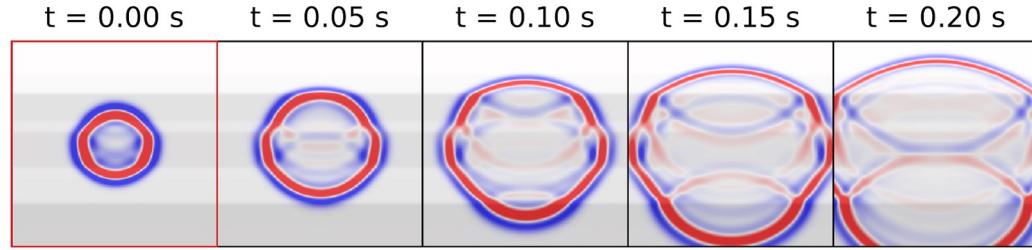
Moseley et al, Solving the wave equation with physics-informed deep learning, ArXiv (2020)

$$L_b(\theta) = \frac{\lambda}{N_b} \sum_j^{N_b} \left( NN(x_j, t_j; \theta) - \underline{u_{FD}(x_j, t_j)} \right)^2 \quad \text{Boundary data from FD simulation (first 0.02 seconds)}$$

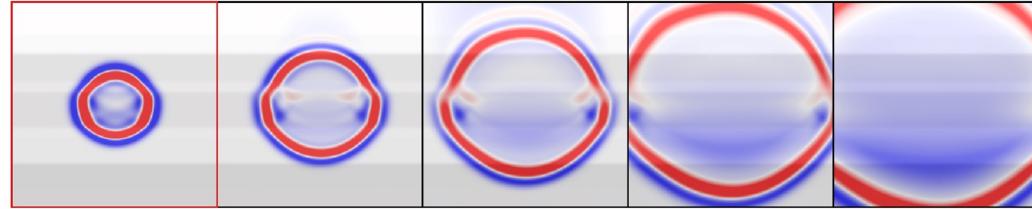
$$L_p(\theta) = \frac{1}{N_p} \sum_i^{N_p} \left( \left[ \nabla^2 - \frac{1}{c(x_i)^2} \frac{\partial^2}{\partial t^2} \right] NN(\underline{x_i}, \underline{t_i}; \theta) \right)^2 \quad \text{Collocation points randomly sampled over entire domain (up to 0.2 seconds)}$$

# PINNs for solving wave equation

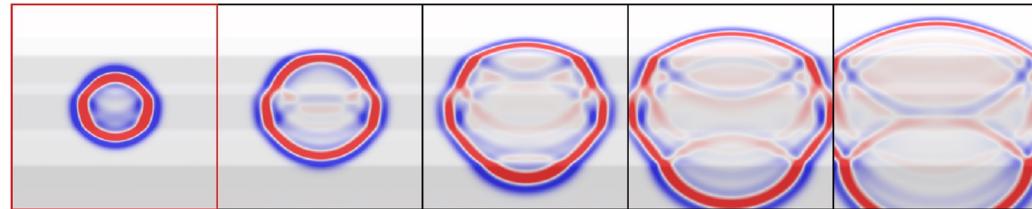
Ground truth FD simulation



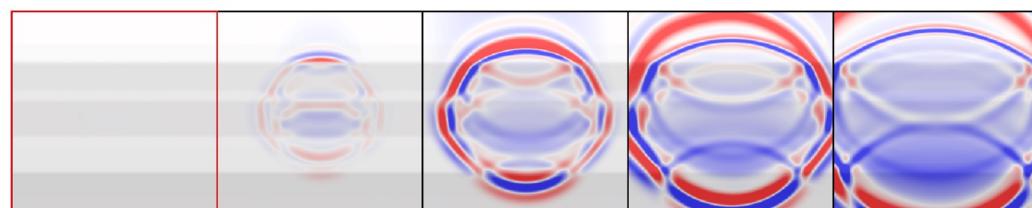
“Naïve” NN



PINN



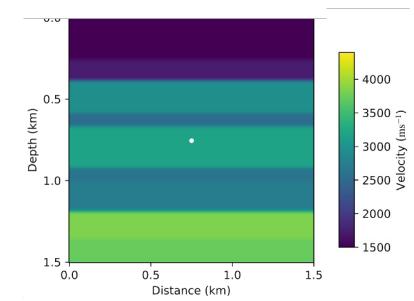
Difference (NN)



Difference (PINN)



Velocity model,  $c(x)$



Moseley et al, Solving the wave equation with physics-informed deep learning, ArXiv (2020)

Mini-batch size  $N_b = N_p = 500$  (random sampling)

Fully connected network with 10 layers, 1024 hidden units

Softplus activation

Adam optimiser

Training time: ~1 hour

# PINNs for inverse problems

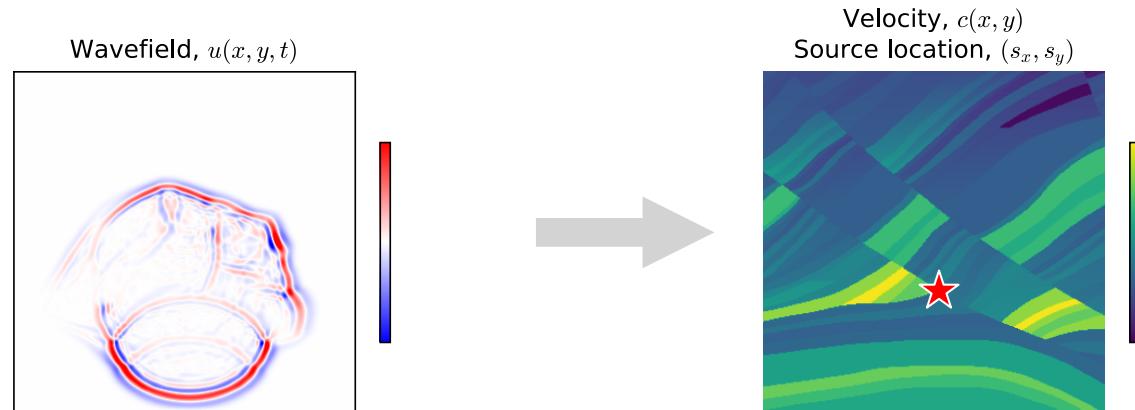
# What is an inverse problem?

Wave equation:

$$\nabla^2 u - \frac{1}{c(x)^2} \frac{\partial^2 u}{\partial t^2} = 0$$

- Fundamentally, inverse problems are **search** problem
- It is often useful to frame them as an optimisation problem, for example:

$$\min_{\hat{a}} \|b - F(\hat{a})\|^2$$



$b$  = observed  
wavefield  $u(x, t)$

$$b = F(a)$$

$a$  = set of input conditions

$F$  = physical model of the system

$b$  = resulting properties given  $F$  and  $a$

# PINNs for inversion

PINNs for solving **forward** simulation:

$$L(\theta) = L_b(\theta) + L_p(\theta)$$

$$L_b(\theta) = \sum_k \frac{\lambda_k}{N_{bk}} \sum_j^{N_{bk}} \| \mathcal{B}_k [NN(x_{kj}; \theta)] - g_k(x_{kj}) \|^2 \text{ Boundary loss}$$

$$L_p(\theta) = \frac{1}{N_p} \sum_i^{N_p} \| \mathcal{D}[NN(x_i; \theta)] - f(x_i) \|^2 \text{ Physics loss}$$

For example:

$$\begin{aligned} D &= \left[ \nabla^2 - \frac{1}{c(x)^2} \frac{\partial^2}{\partial t^2} \right] \\ f &= 0 \end{aligned}$$

# PINNs for inversion

PINNs for solving **forward** simulation:

$$L(\theta) = L_b(\theta) + L_p(\theta)$$

$$L_b(\theta) = \sum_k \frac{\lambda_k}{N_{bk}} \sum_j^{N_{bk}} \|B_k[NN(x_{kj}; \theta)] - g_k(x_{kj})\|^2 \text{ Boundary loss}$$

$$L_p(\theta) = \frac{1}{N_p} \sum_i^{N_p} \|\mathcal{D}[NN(x_i; \theta)] - f(x_i)\|^2 \text{ Physics loss}$$

For example:

$$\begin{aligned} D &= \left[ \nabla^2 - \frac{1}{c(x)^2} \frac{\partial^2}{\partial t^2} \right] \\ f &= 0 \end{aligned}$$

PINNs for solving **inverse** problems:

$$L(\theta, \phi) = L_p(\theta, \phi) + L_d(\theta)$$

$$\begin{aligned} L_p(\theta, \phi) &= \frac{1}{N_p} \sum_i^{N_p} \|\mathcal{D}[NN(x_i; \theta); \phi] - f(x_i)\|^2 \text{ Physics loss} \\ L_d(\theta) &= \frac{\lambda}{N_d} \sum_l^{N_d} \|NN(\underline{x}_l; \theta) - \underline{u}_l\|^2 \text{ Data loss} \end{aligned}$$

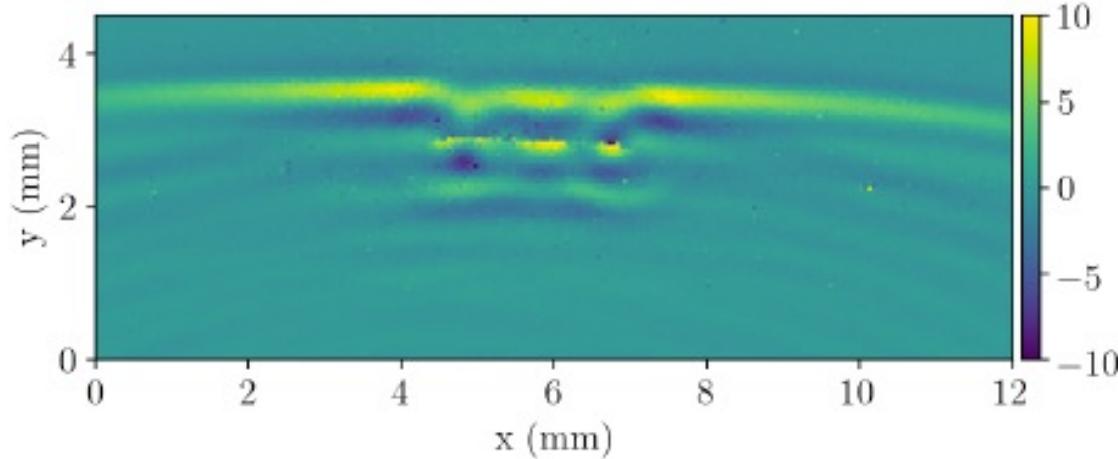
Where  $\phi$  are unknown, underlying PDE parameters we wish to invert for, and  $\{\underline{x}_l, \underline{u}_l\}$  are a set of (potentially noisy) observational data



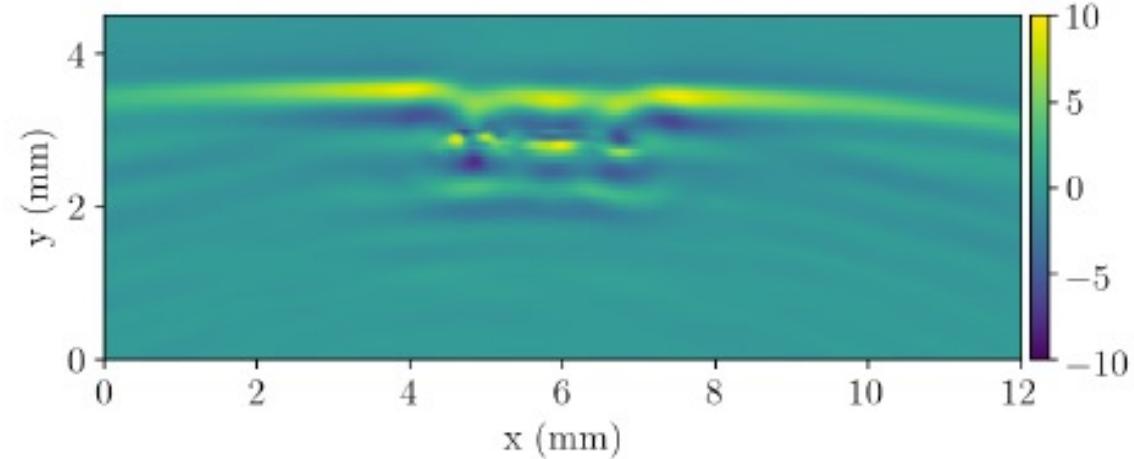
We **simultaneously** learn  $\theta$  and  $\phi$  when training the PINN

# PINNs for wave equation

Shukla et al, Physics-Informed Neural Network for Ultrasound Nondestructive Quantification of Surface Breaking Cracks, Journal of Nondestructive Evaluation (2020)



(a) Actual data at  $t = 12.38 \mu\text{s}$ .

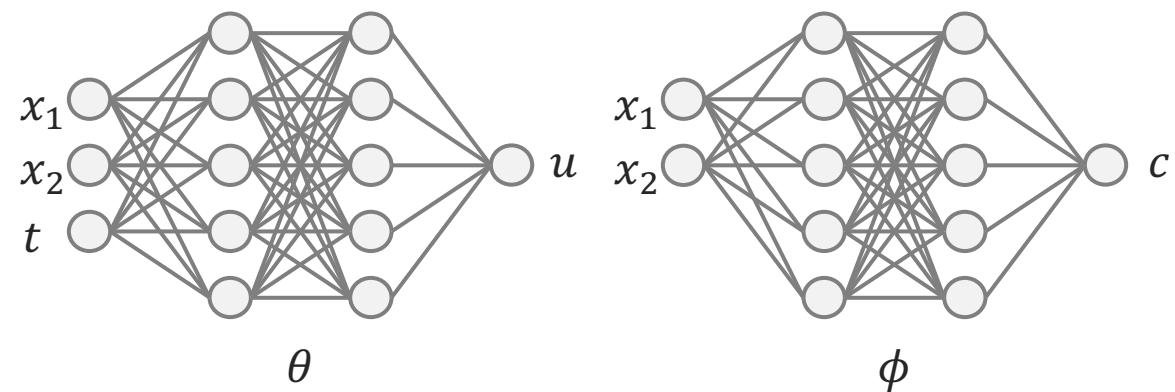


(b) Data recovered from PINN simulation at  $t = 12.38 \mu\text{s}$ .

$$L_p(\theta, \phi) = \frac{1}{N_p} \sum_i^{N_p} \left( \left[ \nabla^2 - \frac{1}{c(x_i; \phi)^2} \frac{\partial^2}{\partial t^2} \right] NN(x_i, t_i; \theta) \right)^2$$

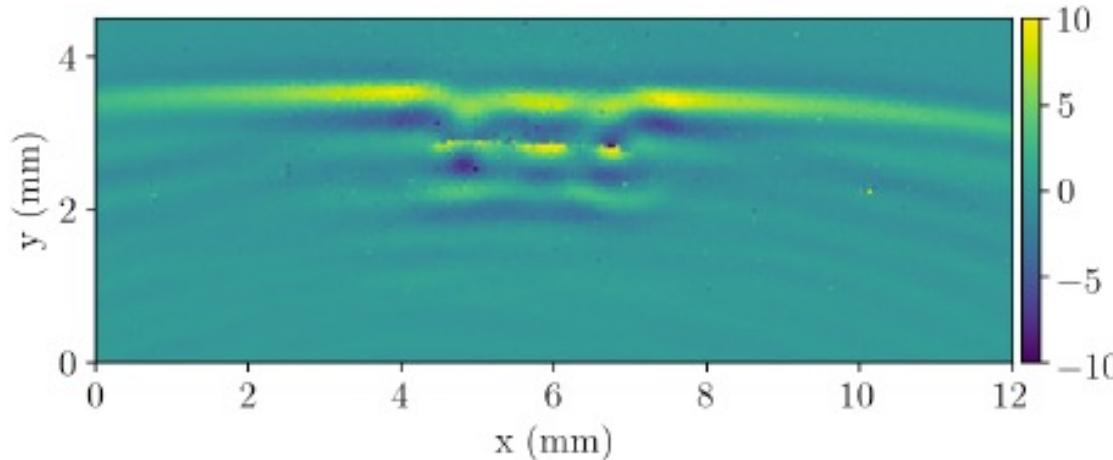
$$L_d(\theta) = \frac{\lambda}{N_d} \sum_l^{N_d} (NN(x_l, t_l; \theta) - u_{obs\ l})^2$$

Treat velocity model as **another** neural network, and simultaneously learn it

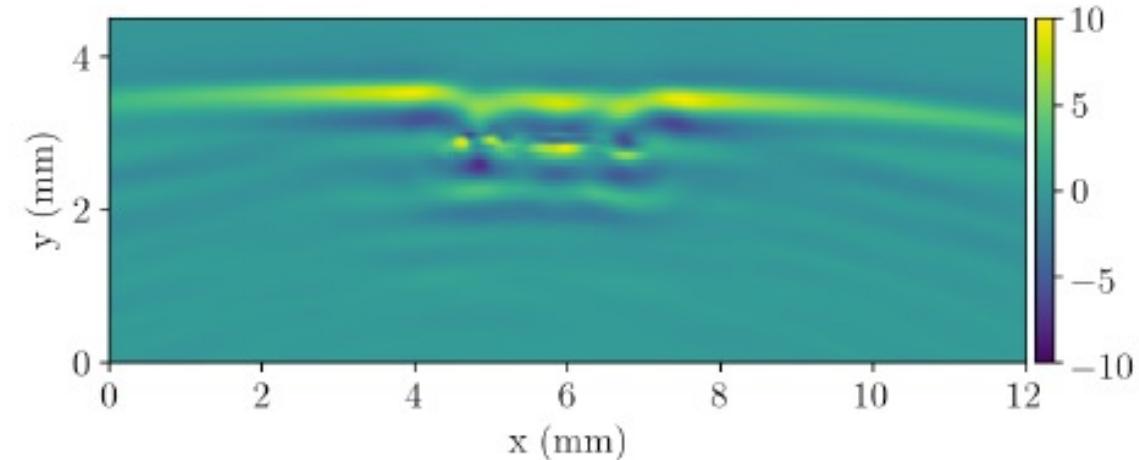


# PINNs for wave equation

Shukla et al, Physics-Informed Neural Network for Ultrasound Nondestructive Quantification of Surface Breaking Cracks, Journal of Nondestructive Evaluation (2020)



(a) Actual data at  $t = 12.38 \mu s$ .

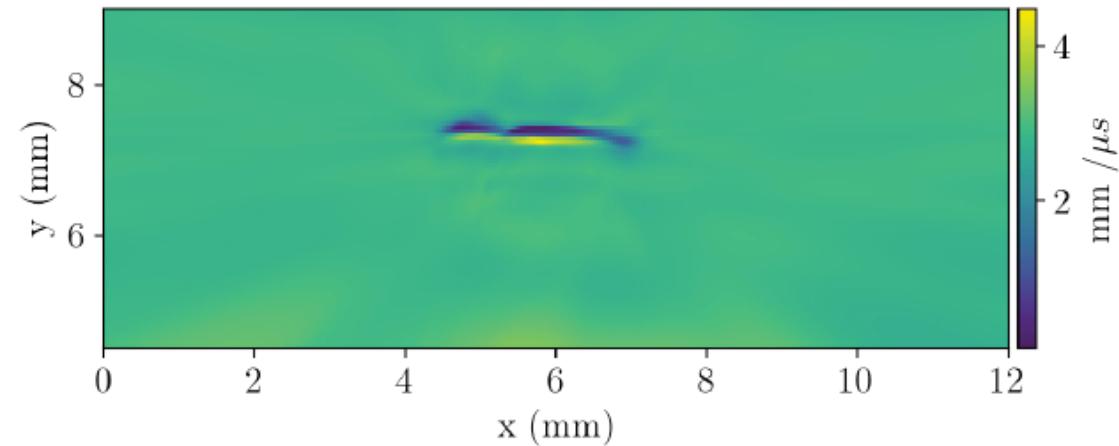


(b) Data recovered from PINN simulation at  $t = 12.38 \mu s$ .

$$L_p(\theta, \phi) = \frac{1}{N_p} \sum_i^{N_p} \left( \left[ \nabla^2 - \frac{1}{c(x_i; \phi)^2} \frac{\partial^2}{\partial t^2} \right] NN(x_i, t_i; \theta) \right)^2$$

$$L_d(\theta) = \frac{\lambda}{N_d} \sum_l^{N_d} (NN(x_l, t_l; \theta) - u_{obs\ l})^2$$

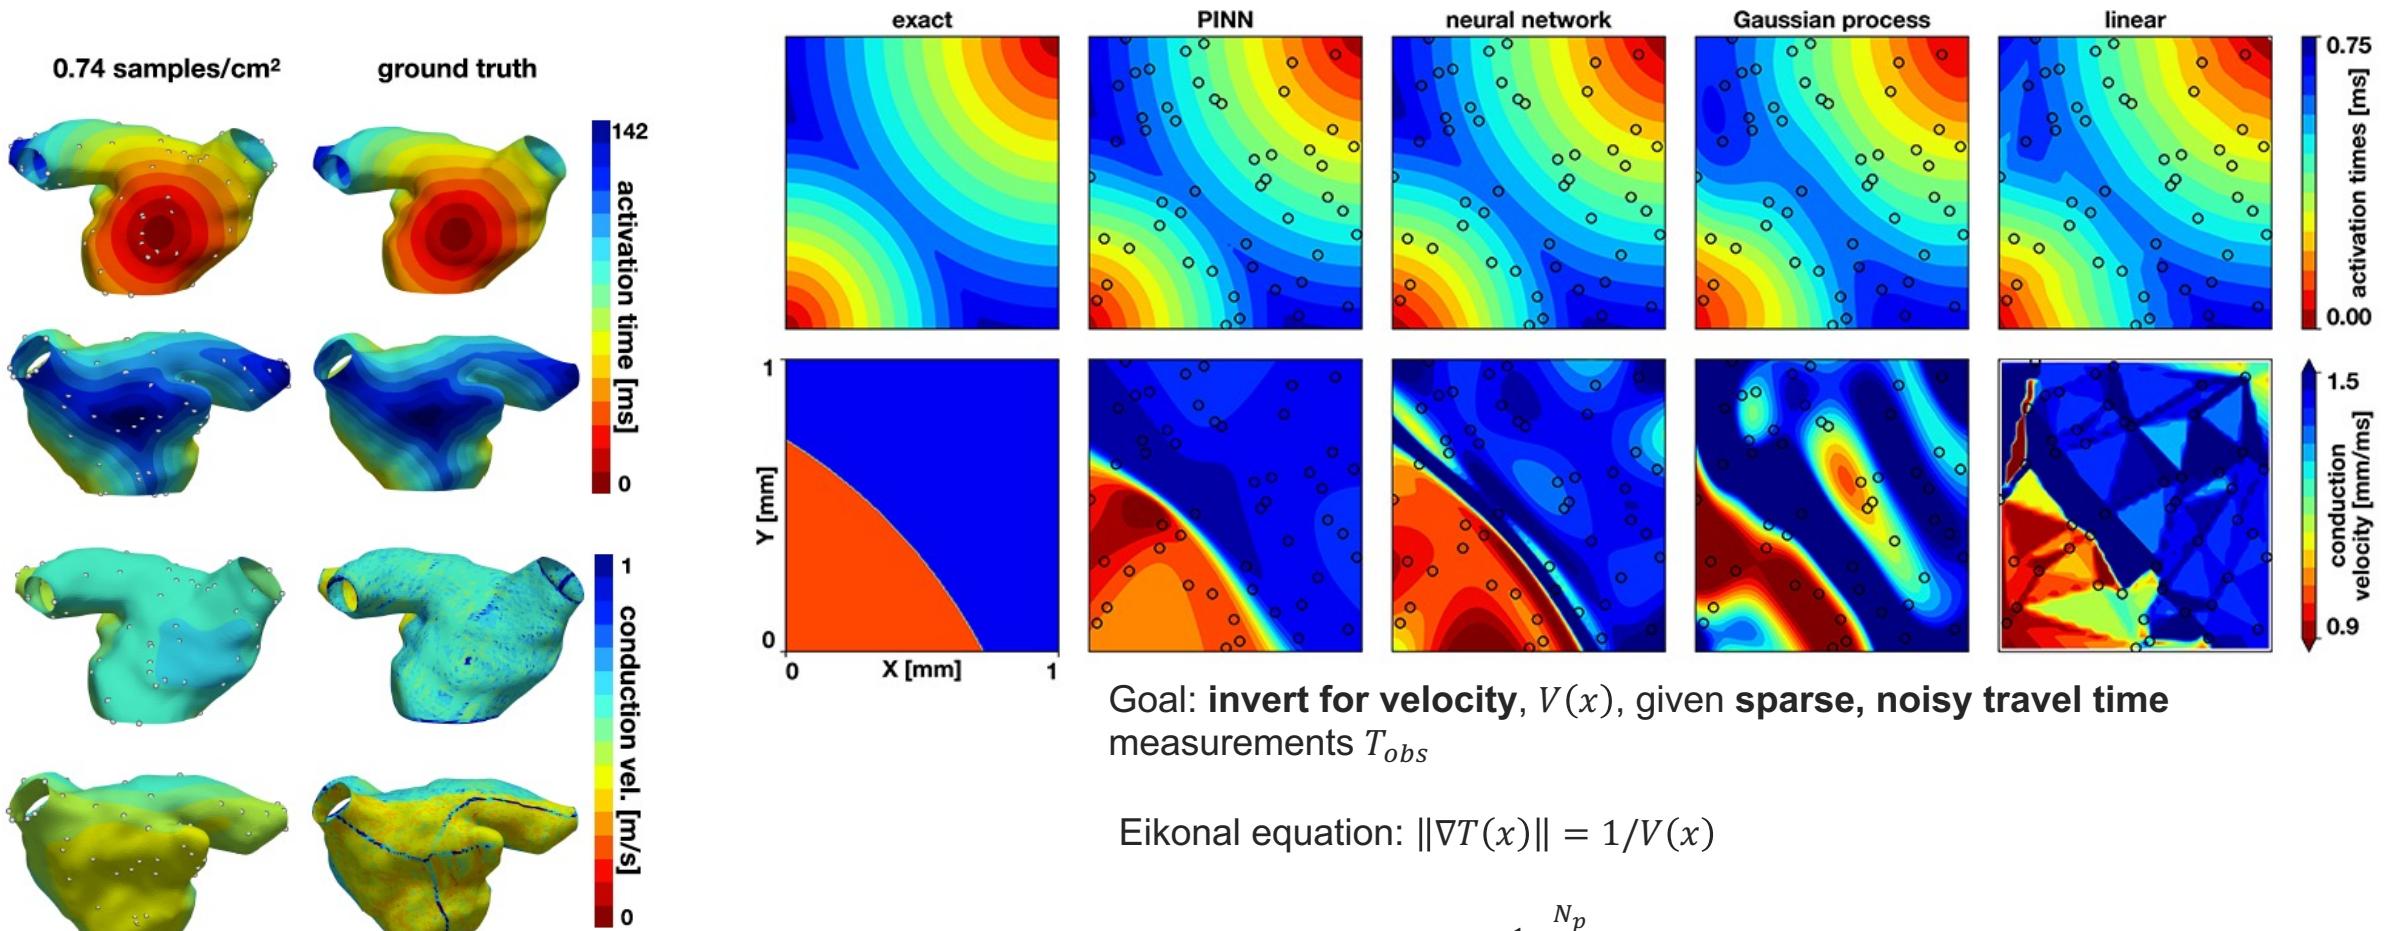
Treat velocity model as **another** neural network, and simultaneously learn it



(d) Speed  $v(x, y)$  recovered from PINN simulation.

# Live-coding a PINN in PyTorch

# PINNs for cardiac activation mapping



Sahli Costabal et al, Physics-Informed Neural Networks for Cardiac Activation Mapping, Frontiers in Physics (2020)

$$L_p(\theta, \phi) = \frac{1}{N_p} \sum_i^{N_p} (\|\nabla NN(x_i; \theta)\| V(x_i; \phi) - 1)^2$$
$$L_d(\theta) = \frac{\lambda}{N_d} \sum_l^{N_d} (NN(x_l; \theta) - T_{obs,l})^2$$

# PINNs for equation discovery

# PINNs for equation discovery

PINNs for solving **inverse** problems:

$$L(\theta, \phi) = L_p(\theta, \phi) + L_d(\theta)$$

$$L_p(\theta, \phi) = \frac{1}{N_p} \sum_i^{N_p} \|\mathcal{D}[NN(x_i; \theta); \phi] - f(x_i)\|^2 \quad \text{Physics loss}$$

$$L_d(\theta) = \frac{\lambda}{N_d} \sum_l^{N_d} \|NN(\mathbf{x}_l; \theta) - \mathbf{u}_l\|^2 \quad \text{Data loss}$$

Where  $\{\mathbf{x}_l, \mathbf{u}_l\}$  are a set of (potentially noisy) observational data

But how do we learn the **entire** differential operator  $\mathcal{D}$ , rather than its parameters  $\phi$ ?

# PINNs for equation discovery

How do we learn an **entire** differential operator  $\mathcal{D}$ ?

Build a **library** of  $n$  operators, such as:

$$\phi = (1, \partial_x, \partial_t, \partial_{xx}, \partial_{tt}, \partial_{xt})^T$$

Then assume the differential operator can be represented as

$$\mathcal{D} = \Lambda \phi$$

Where  $\Lambda$  is a (sparse) matrix of shape  $(d_u, n)$

E.g. for 1D damped harmonic oscillator:

$$\begin{aligned}\mathcal{D} &= (\mathbf{k} \quad \boldsymbol{\mu} \quad \mathbf{m} \quad \mathbf{0}) \begin{pmatrix} 1 \\ d_t \\ d_{tt} \\ d_{ttt} \end{pmatrix} \\ &= \mathbf{m} \frac{d^2}{dt^2} + \boldsymbol{\mu} \frac{d}{dt} + \mathbf{k}\end{aligned}$$

# PINNs for equation discovery

How do we learn an **entire** differential operator  $\mathcal{D}$ ?

Build a **library** of  $n$  operators, such as:

$$\phi = (1, \partial_x, \partial_t, \partial_{xx}, \partial_{tt}, \partial_{xt})^T$$

Then assume the differential operator can be represented as

$$\mathcal{D} = \Lambda \phi$$

Where  $\Lambda$  is a (sparse) matrix of shape  $(d_u, n)$

E.g. for 1D damped harmonic oscillator:

$$\begin{aligned}\mathcal{D} &= (\mathbf{k} \quad \boldsymbol{\mu} \quad \mathbf{m} \quad \mathbf{0}) \begin{pmatrix} 1 \\ d_t \\ d_{tt} \\ d_{ttt} \end{pmatrix} \\ &= \mathbf{m} \frac{d^2}{dt^2} + \boldsymbol{\mu} \frac{d}{dt} + \mathbf{k}\end{aligned}$$

PINNs for **equation discovery**:

$$L(\theta, \Lambda) = L_p(\theta, \Lambda) + L_d(\theta)$$

$$L_p(\theta, \Lambda) = \frac{1}{N_p} \sum_i^{N_p} \|\Lambda \phi[NN(x_i; \theta)]\|^2 + \|\Lambda\|^2 \quad \text{Physics loss}$$

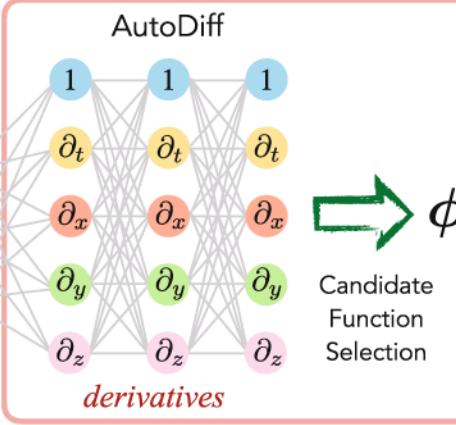
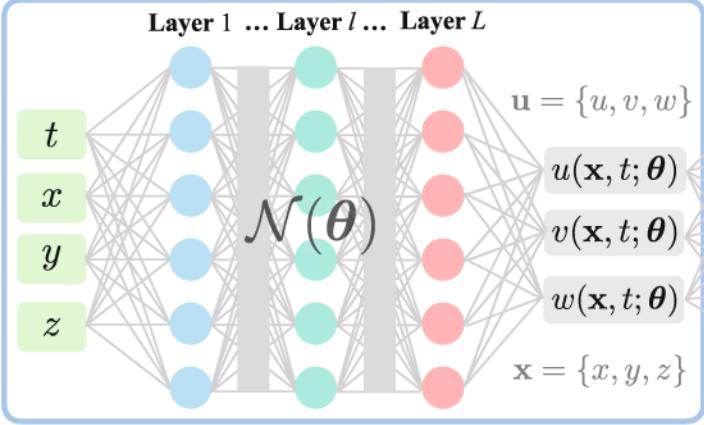
$$L_d(\theta) = \frac{\lambda}{N_d} \sum_l^{N_d} \|NN(\mathbf{x}_l; \theta) - \mathbf{u}_l\|^2 \quad \text{Data loss}$$

Where  $\Lambda$  are treated as **learnable** parameters and  $\{\mathbf{x}_l, \mathbf{u}_l\}$  are a set of (potentially noisy) observational data

Typically, some regularization / prior on  $\Lambda$  (e.g. sparsity) is needed, as this optimisation problem can be very **ill-posed**

# PINNs for equation discovery

DNN with Unknown Parameters  $\theta$



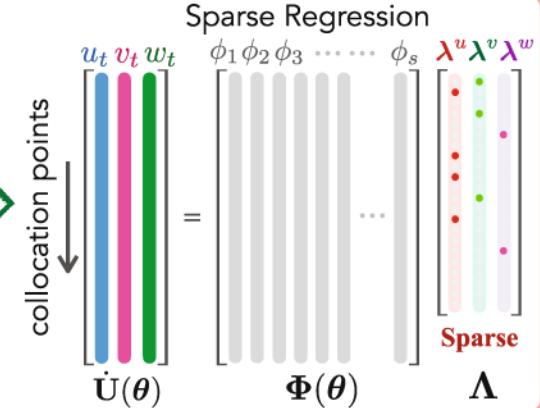
Physical Law with Unknown Parameters  $\Lambda$

PDE Construction

$$\mathcal{R} : \mathbf{u}_t - \phi \Lambda \rightarrow \mathbf{0}$$

where

$$\begin{aligned} \mathbf{u}_t &= \{u_t, v_t, w_t\} \\ \Lambda &= \begin{bmatrix} \lambda_1^u & \lambda_1^v & \lambda_1^w \\ \lambda_2^u & \lambda_2^v & \lambda_2^w \\ \vdots & \vdots & \vdots \\ \lambda_{s-1}^u & \lambda_{s-1}^v & \lambda_{s-1}^w \\ \lambda_s^u & \lambda_s^v & \lambda_s^w \end{bmatrix}_{s \times 3} \\ \text{Note: } \Lambda &\text{ is sparse} \end{aligned}$$



$$\text{Data Loss: } \underbrace{\mathcal{L}_d(\theta; \mathcal{D}_u)}_{\text{measurement}} = \frac{1}{N_m} \|\mathbf{u}^\theta - \mathbf{u}^m\|_2^2 \rightarrow \underbrace{\mathcal{L}(\theta, \Lambda; \mathcal{D}_u, \mathcal{D}_c)}_{\text{total loss}} = \underbrace{\mathcal{L}_d(\theta; \mathcal{D}_u)}_{\text{data loss}} + \underbrace{\alpha \mathcal{L}_p(\theta, \Lambda; \mathcal{D}_c)}_{\text{physics loss}} + \underbrace{\beta \|\Lambda\|_0}_{\text{regularization}} \leftarrow \text{Residual Loss: } \underbrace{\mathcal{L}_p(\theta, \Lambda; \mathcal{D}_c)}_{\text{collocation points}} = \frac{1}{N_c} \|\mathbf{U}(\theta) - \Phi(\theta)\Lambda\|$$

$$\text{Solution by ADO: } \hat{\Lambda}_{k+1} := \arg \min_{\Lambda} \left[ \|\dot{\mathbf{U}}(\hat{\theta}_k) - \Phi(\hat{\theta}_k)\Lambda\|_2^2 + \beta \|\Lambda\|_0 \right] \text{ by STRidge}$$

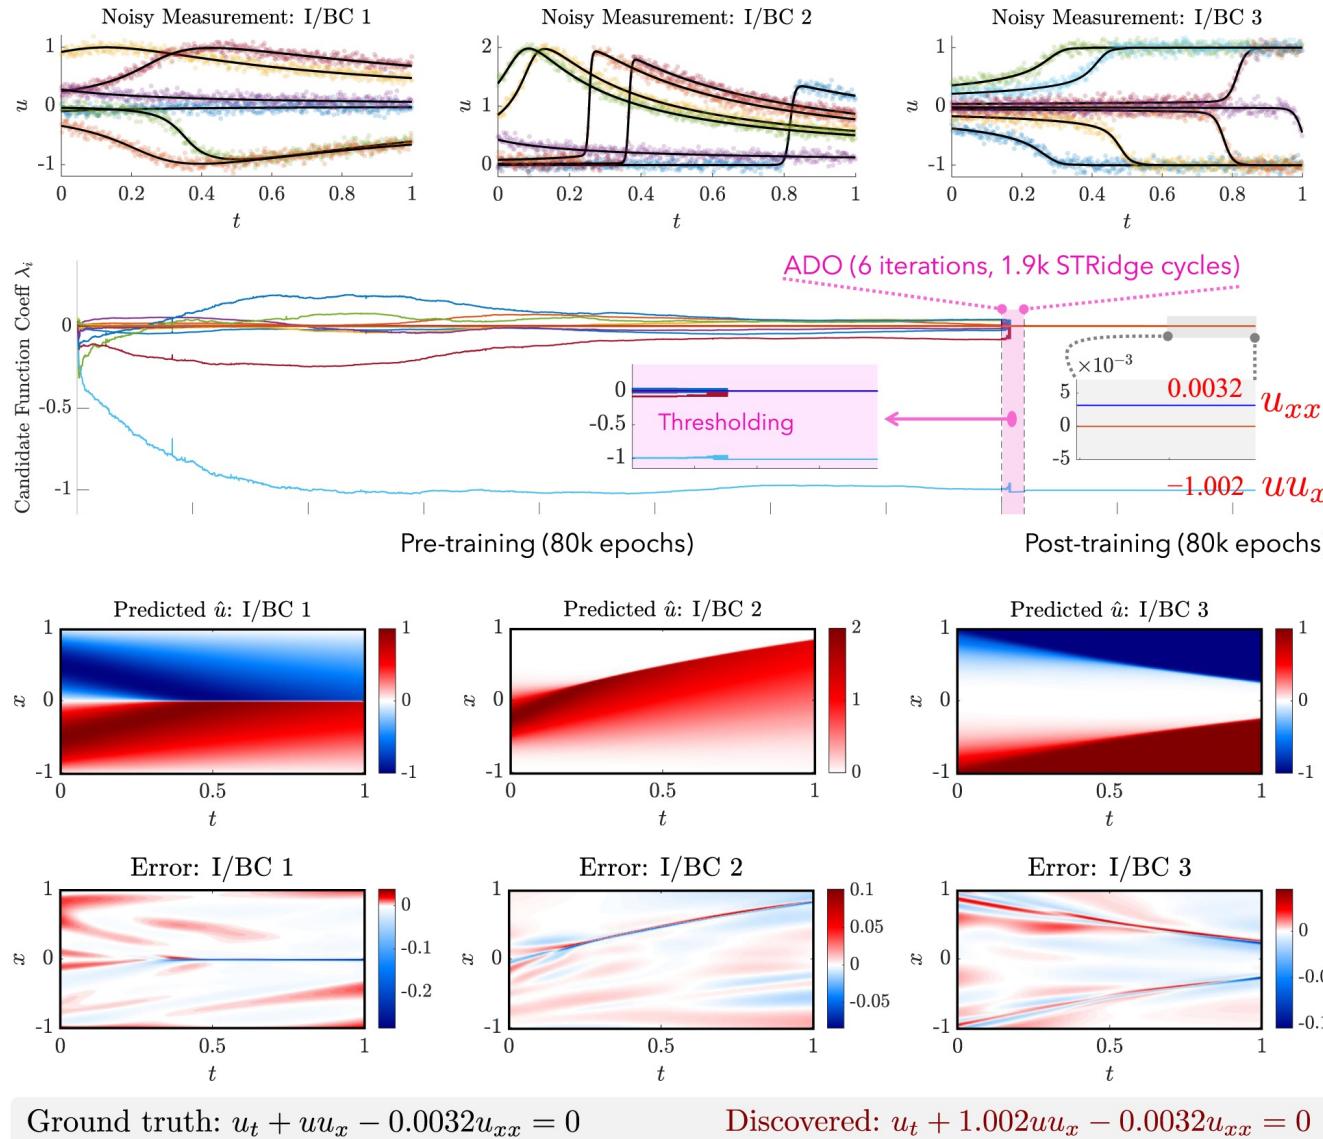
$$\hat{\theta}_{k+1} := \arg \min_{\theta} [\mathcal{L}_d(\theta; \mathcal{D}_u) + \alpha \mathcal{L}_p(\theta, \hat{\Lambda}_{k+1}; \mathcal{D}_c)] \text{ by DNN training}$$



- Trains by alternating between updating  $\Lambda$  and  $\theta$

Chen et al, Physics-informed learning of governing equations from scarce data, Nature communications (2021)

# PINNs for equation discovery



- PINN “discovering” Burgers’ equation
- By combining datasets sampled under three different I/BCs with 10% noise

Chen et al, Physics-informed learning of governing equations from scarce data, Nature communications (2021)

# Lecture summary

- PINNs are an alternative approach for solving **forward, inverse and equation discovery** problems related to differential equations
- They are **mesh-free**
- They use a neural network to **directly approximate** the solution to the PDE